Taming Aspects with Managed Data In Java

Master's Project in Software Engineering

Theologos A. Zacharopoulos

theol.zacharopoulos@gmail.com

May 2016, 83 pages

UNIVERSITEIT VAN AMSTERDAM
FACULTEIT DER NATUURWETENSCHAPPEN, WISKUNDE EN INFORMATICA
MASTER SOFTWARE ENGINEERING
http://www.software-engineering-amsterdam.nl

Contents

List of Abbreviations

AOM Adaptive Object Model.

AOP Aspect Oriented Programming.

CBC Coupling Between Components.

CCC Cross Cutting Concerns.

CDC Concern Diffusion over Components.

CDLOC Concern Diffusion over LOC.

CDO Concern Diffusion over Operations.

DIT Depth of Inheritance Tree.

DSL Domain Specific Language.

JVM Java Virtual Machine.

LCOM Lack of Cohesion in Methods.

LCOO Lack of Cohesion in Operations.

LOC Lines Of Code.

MDSD Model Driven Software Development.

MPO Meta-Object Protocol.

NCLOC Non-Comment Lines Of Code.

NOA Number of Attributes.

OOP Object Oriented Programming.

PP Procedural Programming.

VS Vocabulary Size.

WMC Weighted Method Complexity.

WOC Weighted Operations per Component.

Abstract

The problem of crosscutting concerns results to code scattering and tangling. This in turn significantly affects the modularity of a system, since the separation of concerns principle is violated. Although Aspect Oriented Programming is created in order to solve this problem, several issues arise, the most important being the coupling between aspects and components. This leads to an evolution paradox in Aspect Oriented Software Development.

In this thesis we implement managed data and use them in order to solve the problem of crosscutting concerns. Managed data is a data abstraction mechanism that allows the programmer to define data in addition to their manipulation mechanisms, something that is hard-coded in the programming languages. Defining crosscutting concerns in the manipulation mechanisms, namely data managers, results to a modular way of controlling aspects of data. Therefore, managed data can be used to solve the crosscutting concerns problem and avoid the coupling problem of Aspect Oriented Programming.

Our work focuses on the implementation of managed data in Java, presenting how managed data can be implemented in a static language. Moreover, we demonstrate that managed data can handle aspects by refactoring an existing program's crosscutting concerns. Finally, we provide an evaluation of our managed data refactoring in relation to an Aspect Oriented refactoring of the same use case application, in order to show how managed data can overcome Aspect Oriented Programming problems.

Chapter 1

Introduction

Cross Cutting Concerns (CCC) is a problem the classic programming techniques can not tackle with sufficiently. This results in scattered and tangled code, which affects the system's modularity and its ease of maintenance and evolution. Since Object Oriented Programming (OOP) and Procedural Programming (PP) techniques can not solve this problem, Aspect Oriented Programming (AOP) was introduced [KLM+97] in order to provide a solution to the problem, by presenting the notion of *aspects*.

AOP results in a modular and *single-responsibility* based design, whose properties must be implemented as *components* (cleanly encapsulated procedure) and *aspects* (not clearly encapsulated procedure), both separate concepts that are combined for the result through an automated process called *weaving*. However, relying on AOP, paradoxically does not improve the evolution of a project even though it provides modularity. The main reason is that it introduces tight coupling between the aspects and the application. As a result, the way to address this problem is to consider of a new sophisticated and expressing crosscut language. CCC could be handled on a higher level of the language such as the data structuring and management mechanisms.

Managed data [LvdSC12] allows the developers to take control of important aspects of data as reusable modules. Using managed data a developer can build *data managers* that handle the fundamental data manipulation primitives that are usually hard-coded in the programming language, by introducing custom data manipulation mechanisms. Managed data have been researched and implemented under the Enso project[1], which is developed in Ruby[2] (a dynamic programming language) using its meta-programming framework. Furthermore, it is considered [LvdSC12] that managed data cannot be fully supported in static languages directly, which makes it more challenging for this thesis since our first task it to implement it in Java. In this thesis we use the Java reflection API in order to implement managed data and focus on specific aspects and design patterns implementations using the data managers concept.

Finally, in order to evaluate the implementation of aspects and how we deal with CCC in managed data, we have reimplemented a part of a well-known use case, the JHotDraw, and evaluated the results on a qualitative process. Additionally, we present our results in relation with the original implementation and its AOP counterpart implementation.

1.1 Initial Study

In their study on managed data, Cook et al. [LvdSC12] presented the main idea of managed data, while using a show case of it in a Ruby implementation. As a use case, they presented the Enso[1] project in order to reuse database management and access control mechanisms across different data definitions.

This thesis is an extension of their work; we implement managed data in Java (a static program-

---

[1]http://enso-lang.org/
[2]https://www.ruby-lang.org/en/

ming language) using the Java reflection API[3] and dynamic proxies[4]. Although proxies in static programming languages can not implement the full range of managed data [LvdSC12], Java provides sufficient meta-data manipulation mechanisms through its Reflection API [FFI04] and dynamic proxies. Additionally, our work focuses on the aspects perspective and it provides a solution to the CCC problem by using managed data and their data managers.

The most famous language implementation of AOP is the one provided by Kiczales et al. called AspectJ [KHH+01]. Although AspectJ has been used by a number of projects, some of them with significant research results [HK02] [MM], it includes all the trade-offs of AOP, which are presented in detail in Section 1.2. In this thesis we show how we use managed data in order to handle aspects and compare our findings with the original JHotDraw and an AspectJ show case, the AJHotDraw.

1.2 Problem statement

1.2.1 Problem Analysis

Predefined data structuring mechanisms

One of the most important characteristics of programming languages is the data structures definition. Different types of data structures can be found on different languages and paradigms including *structures*, *objects*, *predefined data structures*, *abstract data types* and more. The common characteristic of these definitions is that they are all predefined. Thus, they do not allow the developers to take control on the data structuring and management mechanisms, but only to create data of these types [LvdSC12].

The problem with this approach is that the predefined data structuring mechanisms can not implement Cross Cutting Concerns and other "common requirements" for data management. In particular, those requirements are not properties that belong to a data structure definition, since, although it is easy to define them individually for every data type, that introduces a significant amount of duplicated and scattered code through the program.

Consequently, in this thesis we implement managed data, which gives the programmers control over the data structuring mechanisms. In addition, it allows the programmer to define the strategies of CCC manipulation as reusable modules.

Crosscutting concerns

As it has be seen [HMK05] there are a number of concerns, during software implementation, that a developer has to work with. For good software modularity, these concerns have to be implemented on different modules, each of these modules implement only one concern [Par72]. However, some of these concerns can not fit to separate modules but their implementation cuts across the system's modules. Those concerns are called Cross Cutting Concerns and result to the problem of *scattered* and *tangled* code.

The problem we study focuses on the CCC that are scattered around an application, resulting in a hard to maintain system by tangling implementation logic and concerns code together. In order to deal with this problem a refactoring of those concerns has to take place, in which the tangled and scattered implementation has to be replaced with an equivalent *aspect* [HMK05].

In this thesis we focus on the modularization of such CCC in aspects, using managed data. We refactor those concerns in modular data structures each of which implement only one concern by lifting the data management up to the application level and allowing the developers to define the concerns in their own data structure manipulation mechanisms.

Aspect Oriented Programming problems

Even though, AOP provides new modularization mechanisms, which should result in easier evolving software, it delivers solutions that are as hard and sometimes even harder to evolve than before

---

[3]https://docs.oracle.com/javase/tutorial/reflect/
[4]https://docs.oracle.com/javase/8/docs/api/java/lang/reflect/Proxy.html

[TBG03]. The problem lays on the aspects, which have to include a crosscut description of all places in the application where this code yields an influence. Thus, the aspects are tightly coupled to the application and this greatly affects the evolvability of the overall system.

Additionally, Steimann [Ste05] argues that modeling languages are not aspect ready. The problem that arises is located at the level of software modeling. More specifically, whereas in OOP roles are tied to the collaborations, in *roles modeling* collaborations rely on interactions of objects and aspects are typically defined independently of one another. Finally, aspects invariably express non-functional requirements, but if the non-functional requirements are not elements of domain models then neither are aspects.

1.2.2 Research Questions

Managed data has not been practically implemented in a static language before, which considered a challenge, therefore our first research questions states *"How to implement managed data in a static language?"*. Based on the argumentation about the relevance of AOP and the solutions that managed data can provide in Cross Cutting Concerns, our second research question is *"Can managed data solve the problem of crosscutting concerns?"*. Finally by using a software showcase, the JHotDraw framework, as well as its AOP implementation AJHotDraw [MM], we evaluate the implementation of managed data on an inventory of aspects. As a result the third research question states *"To what extent can managed data handle an inventory of aspects in the JHotDraw framework, compared to the original and the Aspect Oriented implementation?"*.

1.2.3 Solution Outline

Our solution consists of an implementation of managed data in Java. In particular, we have implemented a framework that can be used in order to implement managed data in Java. This framework provides all the mechanisms of managed data using Java reflection and dynamic proxies. Additionally, one can use the framework in order to refactor the CCC of an existing application.

More specifically we define *schemas* using Java interfaces and dynamic proxies for the *data managers*. Furthermore, we provide as a proof of concept an example given in Enso papers [LvdSC12], but this time developed in Java using our framework. In order to see if managed data solves the problems that AOP introduces, we have implemented an inventory of the following aspects from JHotDraw using data managers:

The Observer Pattern, which as presented in literature [TBG03] [HMK05] [MMvD05a], is by nature not modularized and it scatters "pattern code" through the participant classes. This pattern is considered as a difficult case because it is used a lot in the original JHotDraw source code but with multiple variations, thus it is difficult to extract an abstract version.

The Undo aspect, which is analyzed extensively [Mar04] and a solution is provided by AJHotDraw. More specifically, this aspect consists of aspect-oriented refactoring of the *Command* pattern with *Undo* actions.

We performed aspect refactoring using data managers, that have modularity as a main characteristic, and is evaluated in a new JHotDraw implementation. We compared those aspects with the original version of JHotDraw, and the aspect version, AJHotDraw. Since our solution is a refactoring of the JHotDraw framework we needed a way to ensure the behavioral equivalence between the original and the refactored solution [Fow09]. To archive that, we used the original JHotDraw test suite that consists of 1218 executable tests in total.

1.2.4 Research Method

In order to answer our research questions we studied the theoretical background, we examined our managed data implementation in Java and we evaluated our implementation in an existing use case system.

Managed data implementation in a static programming language. In order to answer the question how to implement managed data in a static language, we have implemented managed data in Java. An extensive presentation of the implementation is given in Chapter 4.

Use case implementation. In order to argue about the contribution of our implementation and managed data in general on solving CCC, we have used a use case application (JHotDraw) which is considered as a good design use case for OOP, along with its AOP implementation (AJHotDraw). Thus, we have built our version of the JHotDraw application (ManagedData-JHotDraw) using our managed data framework to refactor the CCC.

Evaluation and Metrics. In order to evaluate our implementation we have collected a number of metrics presented extensively in Chapter 6. In addition, we present a qualitative comparison between ManagedDataJHotDraw and AJHotDraw based on the modularity properties proposed by Hannemann et al. [HMK05].

1.3   Contributions

Contribution 1: Managed Data Implementation in Java. Our first contribution is the implementation of managed data in a static language, in our case we use Java. The reason we chose Java as the programming language is because Java is a very popular, static, object oriented programming language, with meta-programming (reflective) capabilities which we took advantage of. Managed data implemented in Java, using interfaces for schema definitions and dynamic proxies for the data managers. The final deliverable is a Java library, in `jar` format, which the developer can use to define managed data and data managers for them. Additionally, the developer can define and implement aspects as reusable modules and introduce them in an application without mixing the business logic with the concern logic. More specifically, the schemas and the data managers have to be defined by the developer, as well as any additional functionality that needs to be integrated to the patterns or roles of the application.

Contribution 2: Managed Data Refactoring of JHotDraw. We implemented a new version of the JHotDraw application using our framework in order to evaluate our refactoring of CCC. More specifically, we focused on the *Observer* pattern, which has been used in multiple parts of JHotDraw and cuts "pattern code" on different modules, as well as the *Undo* concern, which is part of a *Command Pattern* and it is scattered through the modules that use this functionality.

Contribution 3: JHotDraw Refactoring Evaluation. Finally, we evaluate our aspect refactoring by using a set of predefined metrics and comparing it in regards to the original version. In addition, we present a set of *modularity properties*, which we discuss in relation to the AOP version.

1.4   Related Work

In this section we discuss the related work of research that inspired this thesis. In particular, we discuss points that we followed and points that we have tried to improve as well as the reason of doing it.

Meta-Object Protocol

The Meta-Object Protocol (MPO) [KDRB91] was first implemented for simple OOP capabilities of the Lisp language in order to satisfy some developer demands including compatibility, extensibility and developers experimentation. The idea was that the languages have been designed to be viewed as black box abstractions without giving the programmers the control over semantics or the implementation of those abstractions. MPO opens up those abstractions to the programmer so he can adjust aspects of the implementation strategy. Providing an open implementation can be advantageous in a wide range of high-level languages and thus MPO technology is a powerful tool for providing that power to the programmer [KDRB91]. Furthermore, MPO provides

flexibility to the programmer because it gives the ability to cleanly integrate something outside the language's scope. Thus, both MPO and managed data allow the programmer to be able to control the interpretation of structure and behavior in a program. However, MPO focuses on behavior of the objects and classes, while in managed data the focus rests solely on the data management. One could conclude that managed data is a subset of the MPO approach since managed data have a more narrow scope.

Adaptive Object Model

Managed data [LvdSC12] is closely related to the Adaptive Object Model (AOM) [YJ02]. AOM is an architectural style that emphasizes flexibility and runtime dynamic configuration. An AOM system, is a system that represents classes, attributes, relationships and behavior as metadata, something that is closely related to the managed data. However, on one hand AOM style is more general than the managed data since it is described at a very high level as a pattern language and it covers business rules and user interfaces, in addition to data management. On the other hand, AOM does not discuss issues of integration with programming languages, the representation of data schemas, or bootstrapping, which are important characteristics of managed data. Finally, AOM is more focused on business systems implementation, not as a general programming or data abstraction technique [LvdSC12].

Model Driven Software Development

Model Driven Software Development (MDSD) refers to a software development method which generates code from defined models. The models represent abstract data that consist of the structure and properties definition of an entity. The idea of the model in MDSD is closely related to the *schemas* in managed data. Similarly to the model definition, schemas define the structure, the properties and any metadata that describe an entity, followed by code generation that adds any extra functionality and manipulation mechanisms to that entity.

The Enso Language

Enso project[5] is the first implementation of managed data, it is open source[6] and is used for EnsoWeb, a web framework written with managed data. Although Enso is implemented in Ruby, which is a dynamic language, the source code was a very helpful resource for our static implementation in Java. The design of Enso was an inspiration for our implementation even though some parts have changed completely in order to conform to our needs and support Java's static type system. Additionally, examples presented in Enso, are also implemented in our case and are presented in Chapter 3.

Aspect Oriented Programming

Although AOP is not directly connected to managed data, it is a mechanism that is relatively easy to be supported in managed data. This mechanism consists of the *weaving* of aspect code in specific join points. The way to support this in managed data is through data managers. How to handle aspects in managed data is one of the topics in this thesis and is going to be presented extensively in the following chapters.

1.5   Document Outline

In this section we outline the structure of this thesis. In Chapter 2 we introduce the background, focusing on the concepts, which the reader must be familiar with in order to follow the next chapters. In Chapter 3 we demonstrate an example to show the capabilities of our implementation. In Chapter 4 the implementation of managed data in Java is presented and discussed, providing detailed explanation of our issues and implementation details. Next, in Chapter 5 a showcase is presented, by applying our implementation to refactor aspects in JHotDraw. Finally, in Chapter 6 an evaluation of the aspect refactoring is illustrated.

---

[5]http://enso-lang.org/
[6]https://github.com/enso-lang/enso

Chapter 2

Background

2.1 Cross Cutting Concerns

There is significant research in software engineering that focuses on the importance of software modularity. The most significant, among the many, advantage of modular systems is their extensibility and evolution [Par72].

However, during the development of a system there is a number of concerns that have to be considered and implemented into the system. In order to follow the modularity principles, those concerns have to be implemented in separate modules, this way the program will be extensible and its evolution will be easier. Nonetheless, many of those concerns can not fit into the existing modular mechanisms in any of the existing programming paradigms including both OOP and PP [KLM⁺97]. In those cases, the concerns are scattered through the modules of the system, resulting to scattered and tangled code. Those concerns are called Cross Cutting Concerns [HMK05]. CCC are considered a serious issue for the evolution of a system and their effects of tangled and scattered code can be disastrous for a system's extensibility.

The reason is that the code which is related to a concern is scattered in multiple modules, while the concern code is tangled with the each module's logic resulting in a system, which does not follow the *Single Responsibility Principle* and consequently is hard to maintain and evolve.

Among many examples of those CCC are persistence, caching, logging, error handling [LL00] and access control. Additionally, some design patterns scatter "design pattern code" through the application, for instance the *Observer Pattern*, *Template Pattern*, *Command Pattern* etc. [HK02] [Mar04].

In order to solve this problem we need a way to refactor and transform the non-modularized CCC into a modular *aspect*. In other words, refactorings of CCC should replace all the scattered and tangled code of a concern with an equivalent module [HMK05], which in AOP terminology is called *aspect* [KLM⁺97].

2.2 Aspect Oriented Programming

Kiczales et al. present AOP [KLM⁺97] by using an example of a very simple image processing application. In that system, as in every system, whenever two properties that are being programmed must compose different tasks and yet be coordinated, they **crosscut** each other. Because general purposes languages provide only one composition mechanism, which leads to complexity and tangling, the programmer must do the co-composition manually. According to their theory, a system's property can be either a *component*, if it can be clearly encapsulated in a generalized procedure, or it is an *aspect*, which is the opposite. AOP supports the programmer in separating components and aspects from each other, by providing mechanisms that make it possible to abstract and compose them together when producing the whole system, a process called *aspect weaving*. Alternatively to the common programming paradigms, OOP and PP, that allow programmers to only separate the *components* from each other but crosscut the concerns.

However, implementing AOP programs is not that easy since several tools are needed in order to

succeed. More specifically, a general purpose language needs only a language, a *compiler* and a *program* written in the language that implements the application. In the case of an AOP based implementation, a program consists of a *component language*, in which the components are programmed, one or more *aspect languages*, in which the aspects are programmed, an aspect *weaver* for the combined languages, a *component program* that implements the components using the component language and finally, one or more *aspect programs* that implement the aspects using the aspect languages. Additionally, an essential function of the aspect weaver is the concept of *join points*, namely the elements of the component language semantics that the aspect programs coordinate with.

2.2.1 AspectJ

There is significant work in the area of aspect oriented languages but one the most important contribution is the AspectJ[1] project, which is a simple and practical aspect-oriented extension to Java and it has been introduced by Kiczales et al. [KHH⁺01]. The authors of AspectJ provide examples of programs developed in AspectJ and show that by using it, CCC can be implemented in clear form, which in other general purpose languages would lead to tangled and scattered code. AspectJ was developed as a compatible extension to Java so that it would facilitate adoption by current Java programmers. The compatibility lays on upward compatibility, platform compatibility, tool compatibility and programmer compatibility. One of the most important characteristics of AspectJ is that it is not a Domain Specific Language (DSL) but a general purpose language that uses Java's static type system. Our goal is to apply those properties for our managed data implementation.

2.2.2 Design Patterns in Aspect Oriented Programming

Hannemann et al. present a showcase of AOP [HK02] in which they conduct an aspect-oriented implementation of the Gang of Four design patterns [Gam95] in AspectJ, in which 17 out of 23 cases show modularity improvements. Even though design patterns offer flexible solutions to common software problems, those patterns involve crosscutting structures between roles and classes / objects. There are several problems that the OOP design patterns introduce in respect to CCC, specifically in cases when one object plays multiple roles, many objects play one role, or an object play roles in multiple patterns [Sul02] (design pattern composition).

The problem lays on the way a design pattern influences the structure of the system and its implementation. Pattern implementations are often binded to the instance of use resulting in their scattering into the code and losing their modularity [HK02].

Even worse, in case of multiple patterns used in a system (pattern overlay and pattern composition), it can become difficult to trace particular instances of a design pattern. Composition creates large clusters of mutually dependent classes[Sul02] and some design patterns explicitly use other patterns in their solution.

2.2.3 Modularity Properties

Hannemann et al. [HK02] provide some example implementations of several design patterns. During the assessment of their findings, the authors used four *modularity properties*. In this thesis we used the same properties to assess our refactoring in relation to the AOP version. The modularity properties are the following:

Locality The *locality* property refers to the ability of an existing class to be incorporated into a pattern instance with effortless adaptation. In this case, all the changes have to be made in the pattern instance.

Reusability *Reusability* holds when a class is not coupled to its role in a pattern. Therefore, it can be used in different contexts without modifications and the reusability of participants can be increased.

---

[1] https://eclipse.org/aspectj/

(Un)pluggability Both the *locality* and the *reusability* properties of a system make the pattern implementations *(un)pluggable*.

Composition transparency Having achieved *locality* and *reusability*, we have obtained an *(un)pluggable* system. This suggests that we can reuse generalized pattern code and localize the code for a particular pattern instance. Thus, creating multiple instances of the same pattern in one application, it is easier to understand *(composition transparency)*. This way the problem of having multiple instances of a design pattern in one application is solved.

2.2.4 Evolvability issues

Since modularization and separation of concerns make the evolution of an application a lot easier and AOP provides mechanisms for modularization and system decomposition, aspect-oriented programs should be easier to be evolved and maintained. Paradoxically, this is not the case since AOP technologies deliver applications that are as hard, and sometimes even harder, as non-AOP.

According to Tourwe et al. [TBG03] the problem is that aspects have to include a crosscut description of all places in the application. Thus, it is much harder to make such crosscuts oblivious to the application and most importantly to the rest of the modules. Additionally, current means for specifying concerns rely heavily in the existing structure of the application, therefore the aspects are tightly coupled to the application (and its structure) and consequently this affects negatively the evolvability of a system since it makes it hard to change its structure. As Tourwe et al. propose, a solution for this problem would be the creation of a new, more sophisticated crosscut language. A language that enables the developer to discriminate between methods based on what they actually do instead of what they look like, in a more intentional way. This new language that implements aspects in a modular way is something try to realize in our thesis.

2.3 Managed Data

Managed data [LvdSC12] is a data abstraction mechanism that allows the programmer to control the definition of the data and their manipulation mechanisms. Additionally it provides a modular way to control aspects of data. Managed data helps the programmers by giving them control over the structuring mechanisms, which until now were predefined by the programming languages. The developers could not take control of them, they could only create data of those types. Managed data provides significant flexibility since it lifts data management up to the application level, by allowing the programmers to build data managers that handle the fundamental data manipulation primitives, normally hard-coded into the programming language.

Managed data has three essential components:

Data description language, that describes the desired structure and properties of data.

Data managers, that enable creation and manipulation of instances of data.

Integration, with a programming language to allow data to be created and manipulated.

In the traditional approach, the programming language includes data definition mechanisms and their processes, which are both predefined. However, with managed data, the data structuring mechanisms are defined by the programmer by interpretation of data definitions.

2.3.1 Schemas

The schemas are the way to describe the structure of the data to be managed. They can be just a simple data description language which programmers can use to describe simple kind of data. For example Cook et al. [LvdSC12] used Ruby hash for the data description on a simple example where the hash was an object that represented a mapping from values to values. However, a simple schema format like this can not be used to describe itself because it is not a record. Therefore, we need a more descriptive language that defines records and fields of records. An instance of a schema is a meta model.

2.3.2 Data Managers

Data managers are the mechanisms that interpret *schemas* with defined manipulation strategies set by the programmers. The input to a data manager is a schema, which describes the structure of the data to be managed. Since the schema is only known dynamically, the data managers must be able to determine the fields of the managed data object dynamically as well. In order to implement such an operation we need a meta-programming mechanism that dynamically analyses the structure of a schema and applies to it the functionality of the data managers. In their implementation Cook et al. [LvdSC12] used the "missing_method" implementation in order to succeed that. In our case we use Java's reflection and dynamic proxies.

2.4 Java Reflection and Dynamic Proxies

The Java programming language provides the programmer with a Reflection API[2] that offers the ability to examine or modify the runtime behavior of applications running in the Java Virtual Machine (JVM). Additionally, Java comes with an implementation of Dynamic Proxies[3] which is a class that implements a list of interfaces specified at runtime.

2.4.1 Reflection

Reflection is the ability of a running program to examine itself and its environment and to change what it does depending on what it finds [FFI04].

In order for this self-examination to be successful, the program needs to have a representation of itself as *metadata*. In a OOP language this metadata are organized into objects, hence they are called *metaobjects*. Finally, the process of the runtime self-examination of these metaobjects is called *introspection*.

Java supports reflection through its reflection API since the version 1.1. Using Java reflection a running program can learn a lot about itself. This information may derive from classes (the `Class` metaobject), class name, class methods, a class super and sub classes, methods (the `Method` metaobject), method name, method parameters, method type, variables, variables handlers and more. Querying information from these metaobjects is called introspection. Additionally to the examining of the these metaobjects, a developer has the ability to dynamically call a method that is discovered at runtime. This process is called *dynamic invocation*. Using dynamic invocation, a `Method` metaobject can be instructed to invoke the method that it represents during the program's runtime.

Although reflection is considered helpful for developing flexible software, it has some known pitfalls:

Security. Since metaobjects give a developer the ability to invoke and change underlined data of the program, it also gives access them to places that are supposed to be secure (e.g. `private` variables).

Code complexity. Consider a program that uses both normal objects and metaobjects. That introduces an extra level of complexity since now a developer has to deal with different kinds of objects on different levels, meta and normal level.

Runtime performance. Of course the runtime dynamic examination and introspection introduce significant overhead on most language implementations. In the case of Java's dynamic proxies a 6.5x overhead observed [MSD15]. However, this is not something that we take into consideration in our implementation.

2.4.2 Dynamic Proxies

Since 1.3 version Java supports the concept of *Dynamic Proxies*. A *proxy* is an object that supports the interface of another object (*target*), so that the proxy can substitute for the target for all practical purposes [FFI04]. A proxy has to have the same interface as the *target* so that it can be used in exactly

---

[2]https://docs.oracle.com/javase/tutorial/reflect/
[3]https://docs.oracle.com/javase/8/docs/api/java/lang/reflect/Proxy.html

the same way. Additionally it *delegates* some or all of the calls that it receives to its target and thus acts as either an intermediary or a substitute object. As a result, a programmer has the capability to add behavior to objects dynamically. The Java reflection API contains a dynamic proxy-creation facility, in `java.lang.reflect.Proxy`.

There are several examples of dynamic proxies implementation in Java including implicit conformance, future invocations [PSH04], dynamic multi dispatch, design by contract or AOP [Eug06].

Proxy Objects

A proxy is an object which conforms to a set of interfaces, for which that proxy was created for. The corresponding proxy class extends the class `Proxy` and implements all its interfaces. Thus, conforming to all those interfaces, a proxy can be casted to any of them and any method defined in those interfaces can be invoked on the proxy object [Eug06].

Invocation Handlers

All the proxy objects have an associated object of type `InvocationHandler`, which handles the method invocations performed on the proxy. Its interface is shown in Listing 2.1.

```java
public interface InvocationHandler {
  public Object invoke(Object proxy, Method method,Object[] args) throws Throwable;
}
```

Listing 2.1: The Invocation Handler Interface

The arguments of the `invoke` method include the object on which the method was originally invoked (i.e., the proxy), the method itself that was invoked on the proxy, and the arguments of that method, if any. Therefore, the `invoke` method is capable of handling any method invocation.

Issues

A proxy instance is an object and it responds to the methods declared by `java.lang.Object`. Thus, when these methods should be invoked and from which object is an issue that arises [FFI04].

The methods `equals`, `hashCode`, and `toString` are inherited by all classes from the `Object` class and they are handled just like custom methods. If they are proxied then they are also overridden by the proxy classes and invocations to them are forwarded to the invocation handler of the proxy. Other methods defined in `Object` are not overridden by proxy classes, as they are `final` [Eug06].

2.5   JHotDraw And AJHotDraw

JHotDraw[4] is a Java GUI framework for technical and structured graphics. It is an open-source, well-designed and flexible drawing framework of around 18,000 non-comment lines of Java code. JHotDraw's design relies heavily on some well-known design patterns [Gam95] and it is considered as a showcase for software quality techniques provided to the OOP community.

The fact that JHotDraw is praised for its design makes it an ideal candidate as a showcase for an aspect oriented migration. Marin and Moonen [MM] use this showcase for adoption of aspect-oriented techniques in existing systems. In particular, they present AJHotDraw[5], which is an aspect-oriented version of JHotDraw developed in Java and AspectJ. The goal of AJHotDraw is to take JHotDraw and migrate it to a functionally equivalent aspect-oriented version.

The authors presented a fan-in analysis of JHotDraw [MVDM04] and implemented an idiom-driven approach to aspect-mining. This way they could extract a number of aspects in JHotDraw. Next, they performed a concern exploration in order to expand their mining results, leading to concern

---

[4]http://www.jhotdraw.org/
[5]http://swerl.tudelft.nl/bin/view/AMR/AJHotDraw

sorts. Concern sorts is a consistent way to address crosscutting concerns in source code [MMVD05b]. This led to the identification and documentation of CCC in JHotDraw, which helps the developers to identify CCC in it. In order to handle the aspects in a more consistent and formal way, Marin et al. provided a list of *template aspect* solutions for their concern sorts. Finally, they performed aspect refactoring of JHotDraw by presenting the AJHotDraw, which according to them, was the largest migration to aspects available to date. Their refactoring aimed at maintaining the conceptual integrity of the original design.

In order to refactor the existing framework, the first thing that AJHotDraw developers needed to do was to create a test subproject for the JHotDraw, called TestJHotDraw[6], which ensures behavioral equivalence between the original and the refactored solution. Refactoring implies preserving the observable behavior of an application [Fow09] and since the developers of AJHotDraw ought to test their functionality, TestJHotDraw was created. There are several contributions of the aspect-oriented implementation approach [MM]. The authors suggest that the project contributes to a gradual and safe adoption of aspect-oriented techniques in existing applications and allows for a better assessment of aspect orientation.

In this thesis we have used JHotDraw and AJHotDraw in order to evaluate our aspect refactoring in managed data. However, TestJHotDraw is written in AspectJ, a language we did not want include in our project, and therefore it is not used. Instead, we used the JHotDraw original test suite, which consists of 1218 test cases, for our refactoring behavioral preservation. The aspect refactoring process is described in Chapter 5.

2.5.1 Refactoring of Crosscutting Concerns

The refactoring of legacy code to aspect oriented code is also known as *Aspect Refactoring* [MMvD05a]. During this process it is important to identify which elements are going to be refactored and which *aspect* solutions will replace them. To evaluate the refactored elements [Fow09], a testing component is needed in order to ensure behavior conservation, hence some coherent criteria to organize CCC are needed. Marin, Moonen and Deursen [MMvD05a] organize the CCC into types, which are descriptions of similar concerns that share the following properties:

- A generic behavioral, design or policy requirement to describe the concern within a formalized, consistent context (e.g., role superimposition to modular units (classes), enforced consistent behavior, etc.),

- An associated legacy implementation idiom in a given (non-aspect oriented) language (e.g., interface implementations, method calls, etc.)

- An associated (desired) aspect language mechanism to support the modularization of the type's concerns (e.g., `pointcut` and `advice`, introduction, composition models).

2.5.2 Role-based Refactoring

The authors present a *role-based* refactoring, which consists of classifying the roles of the pattern in different aspects. The role-based refactoring approach helps the developer to transform a scattered implementation of CCC into an equivalent but modular AOP implementation. Both CCC and refactoring are described in terms of roles.

According to the authors [HMK05], the steps of role-based refactoring are the following:

Selecting a CCC refactoring: The refactoring includes an abstract description of the CCC it targets and a set of instructions to produce a modular AOP implementation of the refactoring (e.g. the Observer pattern CCC refactoring).

Stating a mapping: Map role elements comprising the CCC description to the program elements of the scattered code (e.g. the Subject and the Observer role to concrete classes)

---

[6]http://swerl.tudelft.nl/bin/view/AMR/TestJHotDraw

Planning the refactoring: make the right choices for specific cases since a CCC refactoring involves modifying several parts of a codebase (e.g. naming).

Execution: transform the code according to the refactoring instructions (e.g. modularizes Observer pattern as a result).

Thus, to refactor CCC it is required a mapping from the abstract CCC description to programming components that explicitly describe the CCC implementation.

2.6 Metrics

There is a number of metrics that are used in empirical studies for software assessment. However, most of those metrics refer to OOP systems [CK94] and their basic abstractions such as class, object, methods and attributes. Since AOP introduces a new abstraction, the *aspect*, the available metrics can not be applied to AOP.

Sant'Anna et al. [SGC+03] proposed a new framework for assessing reusability and maintainability qualities of AOP solutions. Those metrics are based on previous work from Chidamber and Kemerer [CK94] and are an extension for measuring aspects as well. Their framework focuses on measurement of the *separation of concerns*, the *coupling*, the *cohesion* and the *size* criteria of an application. The goal of the framework is that software engineers can use it in order to assess design decisions in both OOP and AOP systems. In this thesis we are going to use this framework to assess and discuss the results of ManagedDataJHotDraw in relation to JHotDraw.

The metrics proposed [SGC+03] are grouped in four subsections (criteria) based on the attributes they measure:

Separation Of Concerns refers to the ability to identify, encapsulate and manipulate those parts of software that are relevant to a particular concern [TOHSJ99].

- Concern Diffusion over Components (CDC), counts the number of primary components whose main purpose is to contribute to the implementation of a concern. This metric counts the degree of concern **scattering** in the level of components [FAG+08].

- Concern Diffusion over Operations (CDO), counts the number of primary operations whose main purpose is to assist the implementation of a concern. Constructors also are counted as operations. This metric counts the degree of concern **scattering** in the level of methods [FAG+08].

- Concern Diffusion over LOC (CDLOC), counts the number of transition points for each concern through the lines of code. To use this metric it is required a *shadowing process* that partitions the code into shadowed areas and non-shadowed areas. The shadowed areas are lines of code that implement a given concern. Transition points are the points in the code where there is a transition from a non-shadowed area to a shadowed area and vice-versa [GSC+03]. Therefore, the higher the CDLOC, the more intermingled is the concern code within the implementation of the components; the lower the CDLOC, the more localized is the concern code. This metrics aims to compute the degree of concern **tangling** [FAG+08].

Coupling is an indication of the strength of interconnections between the components in a system. Highly coupled systems have strong interconnections, with program units dependent on each other [Som04].

- Coupling Between Components (CBC), is defined for a component (class or aspect) as the number of other components to which it is coupled.

- Depth of Inheritance Tree (DIT), counts how far down the inheritance hierarchy a class (or aspect) is declared.

Cohesion of a component is a measure of the closeness of the relationship between its internal components [Som04].

- Lack of Cohesion in Operations (LCOO), measures the lack of cohesion of a component. The metric is an extension of Chidamber and Kemerer's [CK94] Lack of Cohesion in Methods (LCOM) metric. More specifically, it measures the amount of method / advice pairs which do not access the same instance variable.

Size measures the length of a software system's design and code.

- Vocabulary Size (VS), counts the number of system components. Those components are classes, in case of OOP, or classes and aspects, in case of AOP. The component instances are not counted.

- Lines Of Code (LOC), counts the number of code lines. Documentation, comments and blanks are omitted.

- Number of Attributes (NOA), counts the internal vocabulary of each component. The metric counts the number of fields of each class or aspect. Inherited attributes are not included in the count.

- Weighted Operations per Component (WOC), measures the complexity of a component in terms of its operations. This metric extends the Chidamber and Kemerer's [CK94] Weighted Method Complexity (WMC) metric for AOP.

Chapter 3

Example Application: State Machine

3.1 Overview

In this chapter in order to show how our managed data implementation works in practice, and in particular in terms of aspect refactoring, we present an example showcase. A detailed demonstration of managed data is presented in Chapter 4; however, in this chapter we focus on its usage.

The example showcase consists of a very simple state machine application. First, we present how to define the schemas of the state machine by using our definition language, the Java interfaces. After that, we show how to define factories that build instances of data, described by our schemas, with interfaces. Next, we define a basic data manager that handles the fundamental functionalities of the state machine's instances. In addition, we also provide a simple program that uses these components in order to describe and interpret the state machine. Finally, having a simple application and a basic data manager, we show that by creating a new data manager the programmer can implement aspects on it. The aspect is implemented in the new data manager, without scattering concern code on the basic application.

A similar example is presented in Enso paper as a showcase for its Object Grammar capabilities [SCL⁺12].

Consider the requirements of the state machine as the following:

- A state `Machine` consists of a number of named `State` declarations.

- Each `State` contains `Transitions` to other states, which are identified by a `name`, when a certain event happens.

- A `Transition` is identified by a certain `event`.

For reasons of simplicity, this example will be a very basic *door* state machine, which includes three states **Open**, **Close** and **Locked**, accompanied by their events: **open_door**, **close_door**, **lock_door** and **unlock_door** respectively. Figure 3.1 illustrates the door state machine.


Figure 3.1: Basic door state machine

To implement this we need to define the models, interpret the definition given from a list of events and finally add any additional functionality (*concern*) needed, in our case we will implement logging of door's current state changes.

3.2   Schemas definition

As a first step, all the models of the state machine program need to be defined. In our implementation we define schemas using Java interfaces with a set of meta data declarations described with Java annotations. For our case, as extracted from the requirements, we need `Machine` (Listing 3.2), `State` (Listing 3.3) and `Transition` (Listing 3.4) schemas.

```java
public interface Machine extends M {
  State start(State... startingState);
  State current(State... currentState);
  Set<State> states(State... states);
}
```

Listing 3.2: The Machine Schema

As it can be seen in Listing 3.2, the `Machine` schema definition interface extends the `M` interface. This describes that this particular interface is a schema definition. The `M` interface has a `schemaKlass` field that is used by the framework to interpret the schema. Section 4.2.2 explains `M`'s existence in detail.

Since we use Java interfaces for our schema definition language, we need a way to define fields and meta data for these fields. For that, we use Java methods for field definitions and Java annotations for meta data on those fields.

As it can be seen in Listing 3.2, the `Machine` schema definition requires a `start`ing state, the `current` state of the machine and a set of `states` that the machine can be into at each time.

Note that for convention, we use a single field (method) definition for both *setter* and *getter*. More specifically, since the method accepts varargs it can be used either without arguments as a *getter*, or with an argument as a *setter*.

```java
public interface State extends M {
  @Key
  String name(String... name);

  @Inverse(other = Machine.class, field = "states")
  Machine machine(Machine... machine);

  List<Transition> out(Transition... transition);

  @Inverse(other = State.class, field = "out")
  List<Transition> in(Transition... transition);
}
```

Listing 3.3: The State Schema

For the `State` schema definition, Listing 3.3, we need a `name` field, which represents the name of the state. This `name` field has been annotated with the `@Key` annotation, which indicates uniqueness when the object is used in a `Set`. The `states` field of Machine can be indexed by name. Moreover, the schema includes a list of `in` and `out` `Transition`s. Next, the field `machine` represents the state machine that the state is part of. As it can be seen in the schema definition, Listing 3.3, the `machine` field has been annotated with `@Inverse`, which indicates that this field is a reference to a field of

another schema. In this case, the `machine` field of `State` schema is a reference to `states` field of `Machine` schema. The same holds for the `in` field, which is inverse of `out`.

```java
public interface Transition extends M {
  String event(String... event);

  @Inverse(other = State.class, field = "out")
  State from(State... from);

  @Inverse(other = State.class, field = "in")
  State to(State... to);
}
```

Listing 3.4: The Transition Schema

Finally, in the `Transition` schema definition, Listing 3.4, we need an `event` that corresponds to the event of the transition. The `from` and `to` fields represent the state that the machine changes from and to respectively. However, these are just references to the `State` schema (Listing 3.3). The *parsing* of schemas is performed through a schema loading process presented in Section 4.2.4.

3.3 Factory definition

Now that we have our schemas, we need a way to build instances of managed objects that these schemas describe. In Java, to create instances of these schemas as managed data, we need to define a factory class, which creates managed data instances (managed objects) for each of these schemas. Note that the method definitions work as `Constructors` of managed objects. Listing 3.5 illustrates a type of this factory for our state machine. This factory class is going to be used by our data manager to (reflectively) parse it and provide an instance of it that creates instances of managed objects.

```java
public interface StateMachineFactory extends IFactory {
  Machine Machine();     // constructor for Machine managed objects
  State State();         // constructor for State managed objects
  Transition Transition();  // constructor for Transition managed objects
}
```

Listing 3.5: The StateMachine Factory

3.4 Basic Data Manager

As mentioned above, in order to interpret and manage the defined schemas of data we need data managers. Our framework includes the definition of a `Basic data manager` that is responsible of interpreting a schema definition to instances of *managed object*s. Conclusively, in order to make a *managed object*, the data manager needs its schema definition (the interfaces that define the schemas) and the factory class (the interface that defines the constructors of the schemas).

3.4.1 A simple program

In the case of a simple program without any concerns, we have to use our schemas to define the state machine, parse it, interpret it and finally use it. A general interpreter for state machines is shown in Listing 3.6.

```
1  void interpretStateMachine(Machine stateMachine, List<String> commands) {
2    stateMachine.current(stateMachine.start());
3    for (String event : commands) {
4      for (Transition trans : stateMachine.current().out()) {
5        if (trans.event().equals(event)) {
6          stateMachine.current(trans.to());
7          break;
8        }
9      }
10   }
11 }
```

Listing 3.6: Door State Machine Interpreter

The definition of the door state machine is shown in Listing 3.7. The `doors` method accepts a `StateMachineFactory`, defines our door state machine and returns the new `Machine` instance.

```
1  Machine doors(StateMachineFactory smFactory) {
2    final Machine doorStateMachine = smFactory.Machine();
3
4    final State openState = stateMachineFactory.State();
5    openState.name(OPEN_STATE); openState.machine(doorStateMachine);
6
7    final State closedState = stateMachineFactory.State();
8    closedState.name(CLOSED_STATE); closedState.machine(doorStateMachine);
9
10   final State lockedState = stateMachineFactory.State();
11   lockedState.name(LOCKED_STATE); lockedState.machine(doorStateMachine);
12
13   final Transition closeTransition = stateMachineFactory.Transition();
14   closeTransition.event(CLOSE_EVENT);
15   closeTransition.from(openState); closeTransition.to(closedState);
16
17   final Transition openTransition = stateMachineFactory.Transition();
18   openTransition.event(OPEN_EVENT);
19   openTransition.from(closedState); openTransition.to(openState);
20
21   final Transition lockTransition = stateMachineFactory.Transition();
22   lockTransition.event(LOCK_EVENT);
23   lockTransition.from(closedState); lockTransition.to(lockedState);
24
25   final Transition unlockTransition = stateMachineFactory.Transition();
26   unlockTransition.event(UNLOCK_EVENT);
27   unlockTransition.from(lockedState); unlockTransition.to(closedState);
28
29   doorStateMachine.start(closedState);
30
31   return doorStateMachine;
32 }
```

Listing 3.7: Door State Machine Definition

In practice, after defining our schemas, the schema loader parses them and provides us with a schema instance. The schema instance can then used by a data manager in order to create a

`StateMachineFactory` instance. A factory instance is created by a basic data manager that provides us with mechanisms that interpret the managed object based on `stateMachineSchema`. Listing 3.8 shows this initialization process.

```java
Schema stateMachineSchema =
  SchemaLoader.load(schemaFactory, Machine.class, State.class, Transition.class);
IDataManager dataManager = new BasicDataManager(); // basic
StateMachineFactory stateMachineFactory =
  dataManager.factory(StateMachineFactory.class, stateMachineSchema);

Machine doorStateMachine = doors(stateMachineFactory);

interpretStateMachine(doorStateMachine, new LinkedList<>(Arrays.asList(
    LOCK_EVENT,
    UNLOCK_EVENT,
    OPEN_EVENT)));
```

Listing 3.8: Door State Machine Initialization

The basic data manager also supports the field accessors of those data, namely, the setters and getters of their values. As it can be seen, the factory is used to create managed objects. The *setup* of the fields is done automatically by the data manager who is responsible for the managed object interpretation.

3.5 Logging crosscutting concern

Now consider a case where we want to add a crosscutting concern at the previous door state machine implementation. A simple concern could be *logging*, which would log every change in the "current" state of the door state machine. Listing 3.9 shows the definition of the logging concern in current states.

```java
public class StateMachineConcerns {
  public static void logCurrentStateChanges(Object obj, String name, Object state) {
    if (name.equals("current")) {
      System.out.println(" > State changed to " + ((State)state).name());
    }
  }
}
```

Listing 3.9: Logging Concern

In order to attach this concern to our machine we need a mechanism that continuously observes the changes (transitions) of the machine's states and reacts accordingly. Usually, this would lead to scattered logging code in the interpretation method or the models themselves (the machine model). This is where data managers come to the rescue. A data manager can implement concerns as modular aspects without scattering code to the components. The programmer can define a manipulation mechanism of his/her data that includes an aspect of preference. Therefore, by implementing our concern with a data manager we can keep the component and aspect code separate.

3.5.1 Observable Data Manager

Regarding the continuous *observation* of our state machine's state changes, we need a data manager that observes these changes in a managed object and executes actions defined by the programmer. More specifically, it has to observe the `Machine`'s current `State` field and perform logging in case

this field's value changes. This data manager creates concrete managed objects as subjects, where observers can be attached in order to be notified of changes and execute an action. It is important to mention that this new data manager has to inherit the basic one in order to include the basic functionality of schema interpretation and field access. This leads to a **stack** of two data managers, each one adding a new aspect of data in a modular way.

During the initialization of our schema factory, we need to use our new data manager, namely `ObservableDataManager`. As Listing 3.10 shows, we change only the data manager object definition in order to generate a new `stateMachineFactory`. The new `stateMachineFactory` instance will create instances that are `Observable`. The `stateMachineFactory` is then passed as an argument to the `door` method in order to build the door state machine. Note that this time, the `Machine` instance is created by using the `Observable` data manager.

The next step it to attach the *logging* concern to our `Machine` object. This is going to be executed in case the `current` state changes.

```java
// Create an observable data manager
IDataManager dataManager = new ObservableDataManager();

StateMachineFactory stateMachineFactory =
    dataManager.factory(StateMachineFactory.class, stateMachineSchema);

// doors method remains the same
Machine doorStateMachine = doors(stateMachineFactory);

// Add logging concern on the state machine
((Observable) doorStateMachine)
    .observe(StateMachineConcerns::logCurrentStateChanges);

interpretStateMachine(doorStateMachine, new LinkedList<>(Arrays.asList(
    LOCK_EVENT,
    UNLOCK_EVENT,
    OPEN_EVENT)));
```

Listing 3.10: Door State Machine Initialization with Observable Data Manager

Concluding, it can be observed that the only parts that have been changed in the original code are the data manager and the logging concern definition. The `door` method remained the same reusing the creation code. The data manager of the `Machine` managed object has been changed to the new observable data manager. Additionally, the logging concern has been attached to the machine object very easily simply by using lambdas.

The first state change happens on `doorStateMachine.start(closedState)` call. After that, by running the program with the commands LOCK_EVENT, UNLOCK_EVENT and OPEN_EVENT, the output is presented in Listing 3.11.

```
> State changed to Closed
> State changed to Locked
> State changed to Closed
> State changed to Open
```

Listing 3.11: Door state machine with logging concern output

The basic data manager allows to solely build managed objects, but the observable data manager also provides the functionality of attaching concerns in the managed objects after a specified event.

The example presented a reusable solution of CCC without scattering and tangling code in the components. The data manager that defines the logging aspect is reusable and modular.

Chapter 4

Managed data in Java

As it has already been mentioned, programming languages include data definition mechanisms that are predefined. This makes them unable to define CCC without repeating and scattering code through the components [LvdSC12]. Notably, the problem is that CCC are not considered features of the data types, but instead features of data management. As a result, we implement managed data to allow the developer to define the mechanisms of data manipulation. This chapter describes our Managed Data implementation in Java, testing our first research question, which states *"How to implement managed data in a static language?"*. It is important to mention that our implementation is inspired by Enso[1], which is written in Ruby. Although Ruby is a dynamic language, Enso significantly contributed to our implementation's design. In this chapter we preset the implementation of managed data in Java, which is available also online as an open-source project called JavaMD (Java Managed Data)[2].

4.1 Overview

In order to build managed data, our framework provides an architecture that has to be followed. The programmer can define *schemas* by using Java interfaces and annotations as the description language. A schema interface has to extend the `M` interface, which determines that this interface is a schema definition. Next, in order to create instances of data of those schemas the programmer has to create *factories* of instances. These factories are implemented in interfaces and define methods that are constructors of data instances. A factory interface has to extend the `IFactory` interface that is simply a marker interface. Finally, to manage this data instances, the programmer has to define *data managers*. The data manager is a two-level proxy in which the programmer has to specify its invocation handlers. The data manager proxy has to extend the `IDataManager` interface. This interface provides a single method, `factory`, which creates instances of *factories*. In the following sections this process will be presented in detail.

In general, to build something we need a schema and a factory. The schema describes the structure of data while the factory creates instances of that data. In order to create a schema, we need to define it by using a description language, in our case Java interfaces. In addition, we need a factory for schemas (schemaFactory), which builds the components of that schema. This is performed by a process called schema loading that returns an instance of a schema. Finally, to manage an instance of data we need data managers. The data managers get a factory definition and a schema instance and returns a new factory instance. This factory instance is used in order to create instances of data described by the schema. Figure 4.1 illustrates this process.

---

[1] https://github.com/enso-lang/enso
[2] https://github.com/TheolZacharopoulos/JavaMD

$$Load : Factory_{schema} \times \overline{Class\langle?\rangle}^T \rightarrow Schema_T$$

$$Manage : Class\langle Factory_T \rangle \times Schema_T \rightarrow Factory_T$$

$$Create : Factory_T \rightarrow T$$

Figure 4.1: Managed data framework

As the figure shows, the framework has main functionalities. First, the schema loader gets a `factory` of schemas [3] and a set of `classes`, which are the interfaces that define schemas. With this information the schema loader parses the arguments, which since they are interfaces it uses Java reflection, and returns a schema instance of `T`. Next, a data manger, gets the schema instance of `T` and the class of the `factory of T` and returns a proxy instance of this `factory`. Finally, by using the `factory of T` the programmer can create instances of `T` based on the behavior described by the data manager.

An example taken from Chapter 3 that uses this process is given in Listing 4.12.

```java
// Load : schemaFactory, interfaces define 'sm' -> sm schema instance
Schema stateMachineSchema =
  SchemaLoader.load(schemaFactory, Machine.class, State.class, Transition.class);

IDataManager dataManager = new BasicDataManager();

// Manage : StateMachineFactory interface, sm schema instance -> sm factory instance
StateMachineFactory stateMachineFactory =
  dataManager.factory(StateMachineFactory.class, stateMachineSchema);

// Create: StateMachineFactory -> Machine
Machine machine = StateMachineFactory.Machine();
```

Listing 4.12: Managed Data Framework Example

4.2 Managed Data Implementation

Managed data allows the programmer to handle the fundamental data manipulation mechanisms using *Data Managers*, one of its distinguishing features being modularity. Using a data description language the programmer defines *Schemas*. Those *Schemas* are parsed by a *Schema Loader* that provide an instance of these schemas. Instance of *Schemas* are the input of *Data Managers*. A *Data Manager* in turn interprets the schema instances that define the structure and the behavior of the data to be managed. *Schemas* and *Data Managers* are essential components of managed data, along with *Integration* in the programming language, in our case being Java.

4.2.1 Data description with Schemas

To create instances of data, we first need to define their structure. *Schemas* describe the outline structure of our data. In order to define *Schemas* in managed data we need a data description language that allows to define records as collections of fields. This language could be anything, e.g. XML, JSON or a different formalism like the one used in Enso. For our implementation we use **Java Interfaces** as a data description language to define records of managed data. By using Java interfaces we use Java's syntax for our definitions. Moreover, Java interfaces use several conventions to encode semantics, for instance Java annotations, which are very useful for meta data definition on *Schemas*.

---

[3] A schemaSchema and its factory are provided by the framework, Section 4.3.1 describes what a schemaSchema is and how it is created.

In addition, this also provides usability to the developer since the IDE can perform a first static type checking.

As a result, to define a *Schema* we first need to define a set of classes that describe that schema. A schema `Klass` [4] is described by a name and a set of `Field`s, each of which has a name and a `Type`. Since Java interfaces are used to define a `(schema)Klass` we need a way to define `Field`s for that `Klass`. A `Field` in our data description language can be defined by using Java's *Method* definition.

Additionally, there are several attributes, considered meta data, that help to define the structure of a `Schema`. In order to define the meta data in our data description language (interfaces), we use *Java Annotations*. Annotations are very declarative in the way they express meta data in interfaces and they are consistent with the system (Java).

Thus, to provide a field with meta data, we define annotations in a *Method* target level since a `Field` is defined by a *Method* declaration Java interfaces.

Note that by using Java interfaces and annotations for our schemas definition, we gain a first level of type checking from JVM. The reason is that before we run our runtime interpretation of schemas, JVM performs type checking in the definitions and in case of wrong types it notifies the programmer. Additionally, this is beneficial when a programmer uses IDE's that perform real time type inspection[5]. In those cases errors on the definitions will be spotted immediately.

The list of the available structure concepts that are supported in our language is presented below [LvdSC12]:

@Key When a method (field definition) is annotated with the `@Key` annotation that determines key when objects of this class are put in a collection (Set). The key should be used on a single field of a Type and its value represents the uniqueness of its Klass's instance.

@Inverse This annotation includes two *annotation element definitions* [6]. The ``Class other'' element defines in which interface the inverse field exist, while the ``String field'' element describes the name of the field that is inverse of. This meta data definition is used as a reference declaration in schemas, meaning that when a programmer updates the value of a field that is annotated with inverse, then the value of the fields that refers to will be also updated. This mechanism is interpreted by the managed object and is used for automated *wiring* of the field across a schema.

@Contain When a field is annotated with the `@Contain` annotation, then this field is considered as *traversal*. In general, traversals describe a minimum spanning tree that is called *spine* and ensures reachability of values. The spine is used in implementations that need a depth-first search by distinguishing between the actual information and the cross-references of the spanning tree. If a spanning tree is defined, then all nodes in a model must be uniquely reachable by following just the spine fields [SCL⁺12]. An example of such functionality is the equivalence between managed objects that is presented in Section 4.5.1. Additionally, traversal fields describe composition, or "is a part of", relationships [LvdSC12].

@Optional When the `@Optional` annotation is on a field's definition this field can include `null` values. `Inverse` fields are `Optional`.

Java Inheritance In addition to the Java annotations, our language uses more Java mechanisms for schemas definition. Java inheritance is one of them. A `schemaKlass` can extend another Klass (super), which works as the traditional Java inheritance, supporting sub typing mechanisms. Implementing this we introduce a *Type Hierarchy* model that includes super and sub classes on managed objects. Note that, since we use interfaces for `schemaKlass`, we implicitly support multiple inheritance because a Java interface can extend more than one interface. However, we do not detect conflicts in case of identical method definitions in the parent interfaces. This could lead to unpredictable functionality and should thus be avoided.

---

[4] We use the "Klass" instead of "Class" convention in order to avoid any kind of ambiguities between Java's Class type and our type system. Klass is used to describe our own class type while Class describes Java's native class type.

[5] https://www.jetbrains.com/help/idea/15.0/code-analysis.html

[6] https://docs.oracle.com/javase/tutorial/java/annotations/declaring.html

Java Collections Finally, another Java mechanism that we use is the definition of a field that includes many values. To define such a field, a programmer has to declare a field's `Type` as a `java.util.List` or a `java.util.Set` of this `Type`.

Using all the aforementioned constructs of our data definition language, a programmer can define any kind of schema, even itself (see Section 4.3). Schema definition examples are presented in Chapter 3 Listings 3.2, 3.3 and 3.4. In those definitions the above concepts can be recognized and their meaning can be revealed in context.

4.2.2 The M interface

In order to declare that an interface is a schema definition, this interface has to extend the `M` interface. The `M` interface defines a single method, namely `schemaKlass`, which refers to the `Klass` that an object is described with. This `Klass` is used by the interpreter (data manager) in order to get the structural description of a managed object. Listing 4.13 illustrates the M interface.

```
public interface M {
  Klass schemaKlass(Klass... schemaKlass);
}
```

Listing 4.13: M Interface

4.2.3 IFactories

However, even if we have the definitions of schemas, we still need a way to create instances of managed data described by them. We can not use Java's mechanisms[7] for this functionality since we need them to be managed data and not ordinary objects. Thus, we use Java interfaces to define instance factories. An `IFactory` interface extension is a list of *constructor definitions* for specific schemas. The `IFactory` interface is an empty interface that is used for documentation as a marker interface (Listing 4.14).

```
public interface IFactory {
}
```

Listing 4.14: IFactory Interface

The methods in the interfaces that extend `IFactory` are used similarly to the constructors in a Java class, while their implementation is handled by the data managers. However, for convention we only define constructors without arguments.

4.2.4 Schema Loading

After defining the schemas using Java interfaces as description language, we need to parse those definitions and create instances of schema.

`SchemaLoader`'s `load` static method takes as input a set of interfaces, which are the schema definitions, and a `SchemaFactory`, which is an `IFactory` that creates `Schema` instances provided by the framework. The `SchemaLoader` returns a new instance of `Schema`. During the reflective analysis (parsing) of the input interfaces, the `SchemaLoader` builds a `Schema` instance that consists of the set of `Types` describing that schema. An example taken from Chapter 3, is shown in Listing 4.15.

---

[7] `new` keyword

```
1  IDataManager dataManager = new BasicDataManager();
2  Schema schemaSchema = new BootSchema(); // Provided by a bootstrapping process
3  SchemaFactory schemaFactory =
4    dataManager.factory(SchemaFactory.class, schemaSchema);
5  Schema stateMachineSchema = SchemaLoader.load(
6    schemaFactory, Machine.class, State.class, Transition.class);
```

Listing 4.15: SchemaLoader Example

As it can be seen, in order to create a `schemaFactory`, the data manager needs a `schemaSchema` instance. The `schemaSchema` is a schema that describes schemas, Section 4.3 explains this further. However, to create instances of a schema we need a `SchemaFactory` but to build a `SchemaFactory` the data manager needs a `schemaSchema` which is again an instance of a schema. This means that we need an existing `schemaSchema` instance before the first use of the `SchemaFactory`, which leads to a circular dependency. This is solved by a bootstrapping process that jump-starts this procedure, Section 4.4 present this in more detail.

In its implementation, the `SchemaLoader` gets as input a `SchemaFactory` and a set of interfaces that describe the state machine schema. Next, it returns a new instance of the state machine Schema. This schema consists of a set of schema Klasses that are described by interfaces, namely `Machine.class`, `State.class` and `Transition.class`. Next, the `SchemaLoader` analyzes the definition of those schemas using reflection and then makes a `Schema` by using the `SchemaFactory` that it has been given. A more detailed description of this process is given in Appendix A. In general `SchemaLoader` can be seen as our parser, which accepts a description of language (the interfaces) and a method create objects of its components (the schema factory).

4.2.5 Data Managers Implementation

However, the schemas are not a complete managed data specification without a corresponding `Data Manager`. A data manager is responsible for interpreting the schema and building virtual objects (managed objects). A managed object's fields are defined by the given schema and acts according to the specifications given by this data manager. Additionally, the data manager ensures that the data given are valid with respect to the schema, including a *type checking* process. More specifically, the data managers describe how a schema definition is handled from the outside world and what its specifications are. These properties may include CCC that can be described separately by special data managers, separating schema and concern definitions. Thus, a managed object can have multiple interpretations based on the data manager that is used to interpret it.

A data manager provides a method `factory` that is used to create `IFactory` instances. Every data manager definition has to implement the `IDataManager` interface, shown in Listing 4.16.

```
1  public interface IDataManager {
2    <T extends IFactory> T factory(
3      Class<T> factoryClass, Schema schema, Class<?>... additionalInterfaces);
4  }
```

Listing 4.16: IDataManager Interface

The `factory` method needs an `IFactory` class and a `Schema` instance. Additionally, it optionally gets a number of interfaces which are added to a proxy instance creation process. This is useful for casting the managed objects on additional types defined by these interfaces. Next, it returns a new `IFactory` instance that creates `Managed Object` instances whose properties are defined by that data manager. The `Schema` argument includes a Set of `Types` (`Primitives` or `Klasses`), while the `IFactory` class declares the constructors of the given schema `Klass`. The data manager through its `factory` method interprets the schemas and builds new `IFactories`, which in turn create `Managed Objects` with the behavior of the data manager.

In the example presented in listing 4.17, Line 2 defines a basic data manager. The data manager gets the `IFactory` class and the `Schema` instance of a state machine as input in the `factory` static method. Next, Line 7 shows a new `IFactory` instance is being created, which builds managed objects with the behavior attached from the basic data manager. Finally, Line 10 illustrates how the managed object instances with this behavior can be built using the `IFactory`.

```
1  // Create a basic data manager for state machines
2  IDataManager basicDataManagerForStateMachines = new BasicDataManager();
3
4  // Create a factory that makes managed objects
5  // with the behavior of the basic data manager.
6  StateMachineFactory stateMachineFactory = basicDataManagerForStateMachines
7      .factory(StateMachineFactory.class, stateMachineSchema);
8
9  // Build an instance of managed object with this behavior.
10 Machine stateMachineInstance = stateMachineFactory.Machine();
```

Listing 4.17: Basic data Manager Example

Basic Data Manager

As described above, we use Java interfaces to define schema Klasses that include fields. Those fields are dynamically discovered by a schema loading process and provided to the data manager as schemas. A data manager has the ability to determine the fields and methods of the managed object during runtime. In addition, when the data manager adds functionality on a managed object then it first delegates the calls to its specifications and then to the fields of an instance. In order to dynamically interpret a schema inside a data manager and delegate its functionalities, we used Dynamic Proxies.

The Two Proxies Process

In our implementation we have separated the Proxy factory (`DataManager`) from the Invocation Handler (`MObject`). This way, the `DataManager` class is responsible for creating proxy instances of `IFactories`. The `IFactory` creates proxy instances with the `MObject` class instances as invocation handlers. The `MObject` instances are responsible for interpreting the schema and delegating actions using their invocation handling implementation. Figure 4.2 illustrates this structure. As it can be seen the data manager is a *factory* that has a single exposed method, `factory()`, that is used to build an `IFactory` instance, which in turn builds `MObject` instances.

The two level *proxying* process that the `DataManager` class performs, from the basic data manager to IFactory and then to MObject, can be seen in Listing 4.18. Note that the `createManagedObject` is `protected` for the sub data managers to override it in order to create MObjects of preference. In addition the `createManagedObjectProxy` is also `protected` in case the programmers wants to to extend the proxy creation.
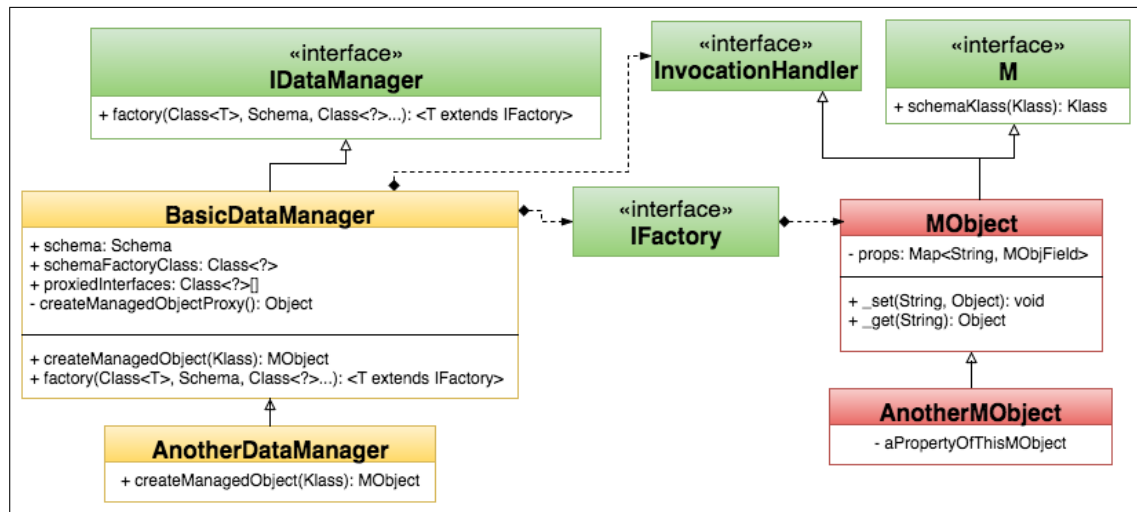
Figure 4.2: Data Manager and MObject

```java
public <T extends IFactory> T factory(
  Class<T> factoryClass, Schema schema, Class<?>... additionalInterfaces) {
  // ...
  return (T) Proxy.newProxyInstance( // First proxy
    factoryClass.getClassLoader(),
    new Class<?>[]{factoryClass},
    (proxy, method, args) -> // inv handler
      createManagedObjectProxy(factoryClass, schema, method, args) // mobject
  );
}
protected Object createManagedObjectProxy(...) {
  //...
  MObject managedObject = createManagedObject(schemaKlass, inits); // inv handler
  return Proxy.newProxyInstance( // Second proxy
    schemaFactoryCallingMethodClassLoader, // schema constructor
    proxiedInterfaces.toArray(new Class[additionalInterfaces.size()]),
    managedObject
  );
}
```

Listing 4.18: Basic Data Manager

Stacking Data Managers

In order to create a stack of data managers that combine their behavior, we use inheritance. Figure 4.2 shows how this works. In detail, `AnotherDataManager` extends `BasicDataManager` and simply overrides the `createManagedObject` and the `factory` methods. The `createManagedObject` method is responsible for creating a new instance of an `MObject`. In this case, the `createManagedObject()` method will create a new `AnotherMObject` instance. The `factory` method is responsible of creating a new `IFactory` instance. Note that it is important that the data managers inherit from a base data manager, leading to the modular aspect of the data managers. As it can be seen, for stacking data managers we used the *Decorator Pattern* [Gam95] which is mentioned also in Cook et al. [LvdSC12] as a strategy for static OOP languages.

4.2.6 MObjects

The `MObject`, is an implementation of the `InvocationHandler` interface. Thus, the `MObject`'s `invoke()` method is called in every field access of the managed object's instance. To manipulate its fields' values this class has two methods, `_set()` and `_get()`. In the implementation of these methods additional checks are performed to ensure the correctness of types and structure of the values. Therefore, a type checker in the schemaKlass level has been implemented in the particular place. The setter and getter methods can be overridden from derived `MObject`s in order to *Decorate* the basic `MObject` with their functionality. Of course they require to call their `supers` for running the type checker.

The `MObject` is the *backing object* that stores a reference to the `schemaKlass` and its implementation represents an instance of that `schemaKlass`. That `schemaKlass` is a meta class that describes the layout of the `MObject` and keeps the `Field`s and their `Types`. During construction, the fields of the `MObject` are specified by its `schemaKlass`. When a field check has to be performed, the `MObject` uses its `schemaKlass`.

Overall, one can easily argue that this class defines the main functionality of managed data. In particular, this class is the *interpreter* of managed data. Therefore, it is responsible for handling the calls to methods (invocation handler), invoke default methods, setup and initialize field values, based on its `schemaKlass`, and internally perform the *type checking*. A detailed presentation of this class and its action is explanation in Appendix B.

4.2.7 Implementing a Data Manager

In order to define a new data manager the programmer needs to define the proxy factory, namely data manager, and its invocation handler, namely MObject. As it can be seen in Figure 4.2, the basic components of a new data manager implementation are the `IDataManager` class (proxy factory) and the `MObject` class (invocation handler).

First, to follow the modularity aspect and the ability to stack data managers together combining their behavior, we need to inherit from, at least, the `BasicDataManager` and its `MObject` respectively. A simple data manager that could be useful is a data manager that introduces immutability to its managed objects. A `Lockable` data manager should first inherit the `BasicDataManager` to get its field access specification. Additionally, it has to add the `Lockable.class` into its `additionalInterfaces` List. This list is responsible of storing interfaces that are going to be added to the proxy instance creation. This way an instance of a proxy can be "casted" to any of these interfaces, in this case the managed object (stateMachine) is casted to `Lockable`. The implementation of the `LockableDataManager` is illustrated in 4.19.

```java
public class LockableDataManager extends BasicDataManager {
  @Override
  public <T extends IFactory> T factory(
    Class<T> factoryClass, Schema schema, Class<?>... additionalInterfaces) {
      // Add the Lockable class in order to use it in the managed object.
      return super.factory(factoryClass, schema,
        pushKlassToProxyInterfaces(Lockable.class, additionalInterfaces));
  }
  @Override
  protected MObject createManagedObject(Klass klass, Object... _inits) {
    return new LockableMObject(klass, _inits);
  }
}
```

Listing 4.19: Lockable Data Manager

Additionally, it should add a *locking* mechanism to ensure immutability of its objects. This is defined in the `Lockable` interface, which is responsible of ensuring the implementation of the specifications.

Listing 4.20 shows the specifications of the interface.

```java
public interface Lockable {
  void lock();
}
```

Listing 4.20: Lockable Interface

Since we have the specifications and the data manager that creates the *Lockable* managed object, we still need the implementation. The implementation is located in the `MObject` and in this case the `LockableMObject`, Listing 4.21.

```java
public class LockableMObject extends MObject implements Lockable {
  private boolean isLocked = false;

  public LockableMObject(Klass schemaKlass, Object... initializers) {
    super(schemaKlass, initializers);
  }

  public void lock() {
    isLocked = true;
  }

  @Override
  public void _set(String name, Object value)
  throws NoSuchFieldError, InvalidFieldValueException, NoKeyFieldException {
    if (isLocked)
        throw new IllegalStateException(
          "Cannot change " + name + " of locked object " + schemaKlass.name() + ".");
    super._set(name, value);
  }
}
```

Listing 4.21: Lockable Managed Object

The `LockableMObject`, by extending the `MObject` and implementing the `Lockable` interface, inherits the basic functionality of a managed object and gets a specification description respectively. Its role is to implement the logic of the immutability, which is as simple as it looks. In order to use this functionality, one needs to create managed objects using this data manager. An example is shown in Listing 4.22.

```
1  IDataManager lockableFactory = new LockableDataManager();
2  PointFactory lockablePointFactory =
3    lockableFactory.factory(PointFactory.class, pointSchema);
4  Point2D lockablePoint = lockablePointFactory.Point2D(1, 2);
5
6  // It was mutable until now, now it is locked (immutable).
7  ((Lockable)lockablePoint).lock();
8  try {
9    lockablePoint.x(2); // Should throw here since its immutable.
10 } catch (IllegalStateException e) {
11   System.err.println("IllegalStateException: " + e.getMessage());
12 }
```

Listing 4.22: Immutability Example

4.3   Self-Describing Schemas

As explained by Cook et al. [LvdSC12], a self-describing schema is a schema that can be used to define schemas, including itself. Our framework is fully self-described, the schemas are also described by schemas which are both models [KBJV06]. To allow schemas to be managed data we need a "self-describing schema mechanism" or *SchemaSchema*. Through the *SchemaSchema* the approach of managed data can be applied at the meta level as well.

The reason that a self-describing schema is important is because schema schemas can be used from factories (IFactory) to create schemas. The schema of schemas is just a schema that allows the creation of schemas, including its own schema [SCL$^+$12]. Additionally, by presenting the schema as the first-class model[KBJV06], they can be extended in the same way just like ordinary models.

4.3.1   SchemaSchema

By using Java interfaces the *Schema* classes are tightly coupled structurally to the Java interfaces used to define them. Since we want to decouple from Java interfaces and reflection we need our own *Klass system*. In order to be self-describing we want this Klass system to be also represented as managed data. To model the structure of a `Schema` itself we need to be able to describe a class as a collection of `Fields`, each of which has a `name` and a `Type` [LvdSC12]. Thus, for our *SchemaSchema* definition we need a `Type`, a `Field` and a `Schema` as a collection of `Types`. A `Type` could be either a `Primitive`, without `Fields`, or a `Klass`, with a set of `Fields`. Additionally, those `Fields` may have some extra meta data attributes that are explained in Section 4.2.1. A schema like this can describe itself since every concept used in the explanation is de facto included in the definition. For a self-describing implementation we need to describe our own SchemaSchema.
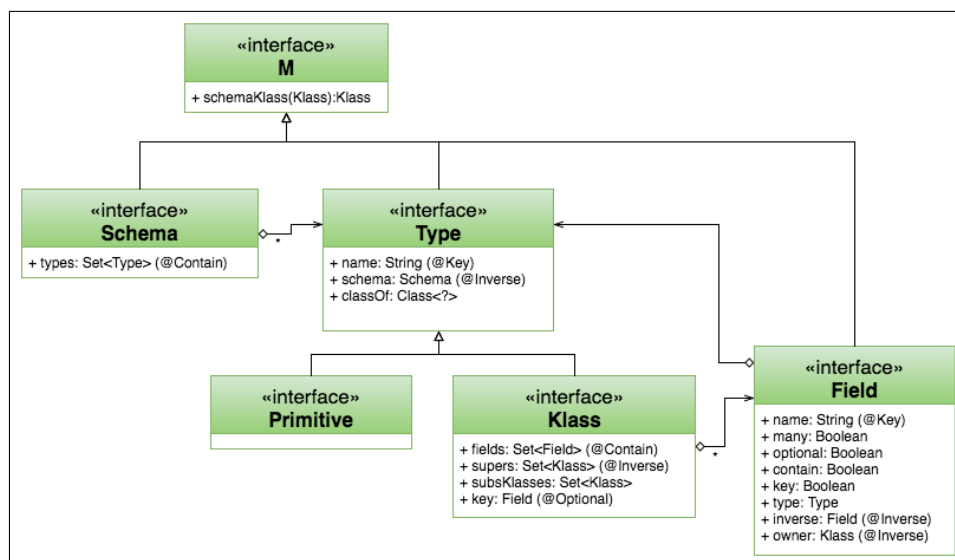
Figure 4.3 illustrates the modeling of this definition.

Figure 4.3: The schema of schemas

4.3.2 SchemaFactory

Considering that we have the schema of our schema (*SchemaSchema*) we need a way to create instances of those *schemaSchema Klasses*. In this case, as we do with the normal schemas, we use an IFactory definition. However, this time it is a *SchemaFactory* that defines constructors of all the schema klasses that are needed to describe our *SchemaSchema*. Listing 4.23 shows its definition.

```
public interface SchemaFactory extends IFactory {
    Schema Schema();
    Primitive Primitive();
    Klass Klass();
    Field Field();
}
```

Listing 4.23: SchemaFactory

To construct the Klass system we need to analyze the Java interfaces using reflection and then build actual instances of the Schema, Klass, Field etc. using the appropriate factory. The `SchemaLoader` is responsible of this process.

4.4 Bootstrapping

Considering that `SchemaSchema` is managed data itself, we can use the `SchemaLoader` to build a new `SchemaSchema`. Nonetheless, we need a description of that `SchemaSchema`, which will be used during the loading process to build the schema Klasses. As a result, we need a *Bootstrap Schema* to jumpstart this process. The *Bootstrap Schema* is exclusively self-describing, as it must manage itself [LvdSC12], and hardcoded in its own class, `BootSchema`.

4.4.1 Cutting the umbilical cord

Having a `BootSchema` in place we can now create "real" `SchemaSchemas` [8]. For consistency, we use those "real" `SchemaSchemas` in order to build other schemas, this way everything is managed data.

---

[8] We call them real because they are managed data and not hard-coded.

After building a real SchemaSchema we no longer need the `BootSchema`, which leads to a process that we call *"Cutting the umbilical cord"*. This process is visualized in Figure 4.4.
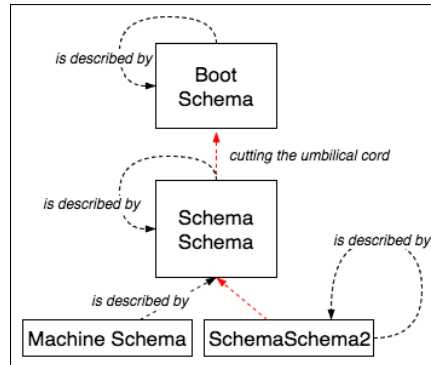

Figure 4.4: Boot Schema models

An example of *"Cutting the umbilical cord"* is shown in Listing 4.4.1, where we use the `BootSchema` to build the `realSchemaSchema` and then we use this `realSchemaSchema` to build another `realSchemaSchema` (`realSchemaSchema2`). In order for a `realSchemaSchema` to be self-described it needs to "cut" its dependency from the previous schemaSchema that describes it. To do that, it has to change its `schemaKlass` to its own `Schema Klass` definition. This can be done by getting the `Schema Klass` instance and setting it to its `schemaKlass`.

```
1  IDataManager dataManager = new BasicDataManager();
2  SchemaFactory schemaFactory =
3    dataManager.factory(SchemaFactory.class, new BootSchema());
4  Schema realSchemaSchema = SchemaLoader.load(
5      schemaFactory,
6      Schema.class, Type.class, Primitive.class, Klass.class, Field.class,
7      Primitives.class);
8
9  Klass realSchemaKlass = // find the "Schema" Klass
10   realSchemaSchema.klasses().stream()
11   .filter(klass -> "Schema".equals(klass.name()))
12   .findFirst().get();
13
14 realSchemaSchema.schemaKlass(realSchemaKlass); // Self describing
15
16 IDataManager dataManager2 = new BasicDataManager();
17 SchemaFactory schemaFactory2 =
18   dataManager2.factory(SchemaFactory.class, realSchemaSchema);
19 final Schema realSchemaSchema2 = SchemaLoader.load(
20       schemaFactory2,
21       Schema.class, Type.class, Primitive.class, Klass.class, Field.class);
22
23 Klass realSchemaKlass2 = // find the "Schema" Klass
24   realSchemaSchema2.klasses().stream()
25   .filter(klass -> "Schema".equals(klass.name()))
26   .findFirst().get();
27
28 realSchemaSchema2.schemaKlass(realSchemaKlass2); // Self describing
```

Listing 4.24: Cutting the umbilical cord

The `BootSchema` is used in order to describe the SchemaSchema, making the SchemaSchema independent and managed data itself. Thus, it can be used to create other schemas like the Machine schema or even itself.

4.4.2 Primitives Definition

Since the Bootstrap Schema defines the primitive types for its description, the real schema schema needs a way to include them as well. These initial Java primitives supported in our implementation are shown in Table 4.1.

	Class	Name	Default Value
Integer	Integer.class	"Integer"	0
int	int.class	"int"	0
Boolean	Boolean.class	"Boolean"	false
boolean	boolean.class	"boolean"	false
String	String.class	"String"	""
Double	Double.class	"Double"	0.
Float	Float.class	"Float"	0.f
Class	Class.class	"Class"	null
Object	Object.class	"Object"	null

Table 4.1: Primitives Table

To define those primitives we use an interface called *Primitives*, introduced during the loading of the real schema, as seen in Line 7 of Listing 4.24. The definition of this interface is shown in Listing 4.4.2 which is a simple `Class/Name` mapping [9].

```
public interface Primitives {
  Integer Integer();
  int _int();
  Boolean Boolean();
  boolean _boolean();
  String String();
  Class Class();
  Float Float();
  Double Double();
}
```

Listing 4.25: Primitives Definition

The benefits of such a definition is that the `Primitives` interface is extensible. By extending it, one can add more primitives in the schema, as long as it is introduced during the schema loading. This extension is very helpful since it allows the programmer to define a number of classes that do not need to be managed data. Consequently, while implementing a new application, the programmer does not need to make everything as managed data. Chapter 5 describes how this can be used in practice.

4.5 Implementation Issues

The fact that we use Java reflection and dynamic proxies, along with the fact that everything is managed data, even the schemaSchema, introduces some issues.

---

[9]We use the "_" prefix convention in order to define names of primitives that are reserved words in Java.

4.5.1 Equivalence

The `bootstrapSchema`, `realSchemaSchema` and `realSchemaSchema2` managed objects from the Listing 4.4.1 should be equal because they ultimately describe the same *Schema*.

However, since, apart from the `bootstrapSchema`, they are managed data and not normal Java objects, we need a way to check for equality on managed objects. We have implemented the equivalence functionality for managed objects, using the *Equality Checking for Trees and Graphs algorithm* by Michael D. Adams and R. Kent Dybvig [AD08].

4.5.2 The classOf field

As it has be presented in Section 2.4.2, for a proxy object to conform with interfaces and be casted to any of them, it needs these interfaces during its initialization. To support that, we have added the `classOf` field in the `Type` schema Klass, which is of type `java.lang.Class` and is a reference of the Java class that this schema Klass is described from. Listing 4.26 shows how this has been used in the basic data manager. The value of the `classOf` field has been set during the schema loading process to the interface that describes the schema klass.

```java
public <T extends IFactory> T factory(
  Class<T> factoryClass, Schema schema, Class<?>... additionalInterfaces)
{
  // Add the classOf interfaces of the schema Klass
  for (Klass klass : schema.klasses()) {
    this.addProxyInterface(klass.classOf());
  }
  //...
}
```

Listing 4.26: Use of classOf on BasicDataManager

4.5.3 Hash-code of Managed Objects

For our framework, the data managers depend on the fields that are annotated with `Key` in order to define uniqueness. We do not depend on the ordinary `hashCode` for managed objects, we do not call it and therefore we have not implemented it. If it is a collection (Set) field type, then the type has to have a `Key` field. In this case, we obtain the value of the key field and index it into a `HashMap`.

Using the `Key` field as the key of the hashmap works whether it is a primitive or not since we get the `Object.hashCode()` of that key. However, that suggests that the key is not of our schema Klass system but a Java type. Finally, the `MObject` invocation handler delegates the call of the `hashCode` method to the real object so that it would never fail, although this is not suggested because it may lead to unpredictable results.

4.5.4 Java 8 Default Methods

Java 8 supports the definition of default methods in interfaces. According to the specification[10], default methods enable the programmer to add new functionalities to the interfaces and can be used as method implementation in abstract classes. We use Java 8 default methods in order to add functionality to our schema definitions. In particular, methods that are defined as *default* are ignored during the parsing and no fields are created for them. We consider this as a helpful mechanism for defining functionality inside the schemas. A notable feature is that the default method invocation in the MObject is `protected`, which makes it possible for the derived data managers to "monitor" when a default method is invoked.

An example of a *computed* field, defined by a default method, is shown in Listing 4.27. The `Point2D` has two fields, x and y, and defines a computed field using a Java 8 default method. The

---

[10]https://docs.oracle.com/javase/tutorial/java/IandI/defaultmethods.html

default method translates this point, at location (x,y), by dx along the x axis and dy along the y axis, so that it now represents the point (x + dx, y + dy).

```java
public interface Point2D extends Point {
  Integer x(Integer... x);
  Integer y(Integer... y);

  default void translate(int dx, int dy) {
    this.x(this.x() + dx);
    this.y(this.y() + dy);
  }
}
```

Listing 4.27: Computable Field

Chapter 5

Aspect refactoring of JHotDraw

5.1   Crosscutting Concerns Identification

Our managed data framework can implement aspects in modular data managers. Yet, to solve the problem of CCC it is first required to identify them in the source code. This leads to a process called *aspect mining*. *Aspect mining* is a reverse engineering process that aims at finding CCC in existing systems [MVDM04]. The aspect mining topic has been addressed in previous research that include methods such as clone detection [BVDVET05], machine learning [SGP04], IDE tools [RM02] and more. Marin et al. [MVDM04] introduced a technique constructed by spotting methods that are invoked from many different places (high fan-in), in order to identify candidate aspects in open-source Java systems. One of these projects include the JHotDraw. In this thesis we focused on their concern findings in refactoring JHotDraw. In particular, we focused on the *FigureSelection*, concern, which is an observer pattern implementation and the "Undo" concern.

5.2   Aspect Refactoring in Managed Data

In order to evaluate the ability of managed data to handle aspects, we have refactored two CCC that have been identified in JHotDraw. More specifically, in this chapter we investigated both the `FigureSelectionListener` and the `ChangeAttributeCommand` (*Undo*) CCC.

First, the `FigureSelectionListener` is an observer pattern implementation that is used to notify a subject for changes on figure selection in `DrawingView` of JHotDraw. In order to implement this concern we use a new data manager that defines a *subject* role on the `DrawingView`. This data manager allows the programmer to attach observers that need to listen to figure selection changes. Second, the *Undo* concern is a command pattern that is used by the a number of components in JHotDraw that supports the undo functionality. For our implementation we use a new data manager that implements the *undoable* functionality on commands. This way the undo concern is defined in the data manager and any command that needs this functionality uses the proper factory.

The choice of those concerns has been made on purpose, since those are the concerns that AJHot-Draw refactors using AspectJ and AOP techniques. For the refactoring we used our implementation of managed data in Java, presented in the previous chapter. Therefore, by having three versions of the same application (JHotDraw) and by solving the same concerns we will be able to perform a comparative evaluation. The three systems included in our assessment are: **JHotDraw**[1], the original OOP version, **AJHotDraw**[2], the AOP refactored version and our **ManagedDataJHotDraw**[3], the managed data refactored version. We focused on those concerns because they were also identified, solved, analyzed and presented in AJHotDraw. Note that, for compatibility and comparison reliance, we used the version *JHotDraw v.5.4b1* since AJHotDraw also refactors the same version.

In order to refactor JHotDraw, we first had to migrate it in managed data. The result of this

---

[1] http://www.jhotdraw.org/
[2] https://sourceforge.net/projects/ajhotdraw/
[3] https://github.com/TheolZacharopoulos/ManagedDataJHotDraw

migration is available on an open-source project, the ManagedDataJHotDraw. We claim that this is the first aspect refactoring of an application using managed data to date, since this project aims on showing how managed data can deal with CCC in existing systems.

5.3 Migration Process

The refactoring of an application of JHotDraw's size required a significant amount of time to study and familiarize with, yet, its well-designed OOP code, made it easy to grasp. We solely focused on the parts that were going to be refactored, based on refactorings that AJHotDraw developers [MM] performed. Thanks to their fan-in analysis [MVDM04], we targeted the same concerns in order to make a fair comparison. Furthermore, during the implementation of ManagedDataJHotDraw we focused on maintaining behavioral coherence and the original design. The migration process of the basic components is detailed described in Appendix C. Something that has to be noted in this process is the `JHotDrawPrimitives` interface, which defines the objects that are not managed data (primitives). This made the migration process a lot simpler, since we only had to implement a part of it, the components that were needed to be managed data.

However, Java attribute keywords such as "synchronized" can not be supported on default methods. Moreover, visibility is an issue since all default methods are `public`, which means that the encapsulation is violated. Finally, private classes definition is not possible inside schemas, although they can be defined outside as managed data.

5.4 Aspect Refactoring of JHotDraw

Aspect refactoring usually refers to the refactoring of legacy code in aspect oriented code. However, in this section we present an aspect refactoring of JHotDraw legacy code in managed data.

5.5 FigureSelectionListener: Observer Pattern

The `FigureSelectionListener` observer pattern of JHotDraw is a concern first presented by Hannemann et al. [HMK05] in their role-based refactoring of design patterns in AspectJ. Later, Marin et al. used the same concern and migrated it into their AJHotDraw implementation [MMvD05a]. Likewise, we have also implemented the same aspect for our refactoring in order to compare our aspect solution with the existing one.

5.5.1 FigureSelectionListener in JHotDraw

The original `FigureSelectionListener` observer pattern of JHotDraw is illustrated in Figure 5.1.
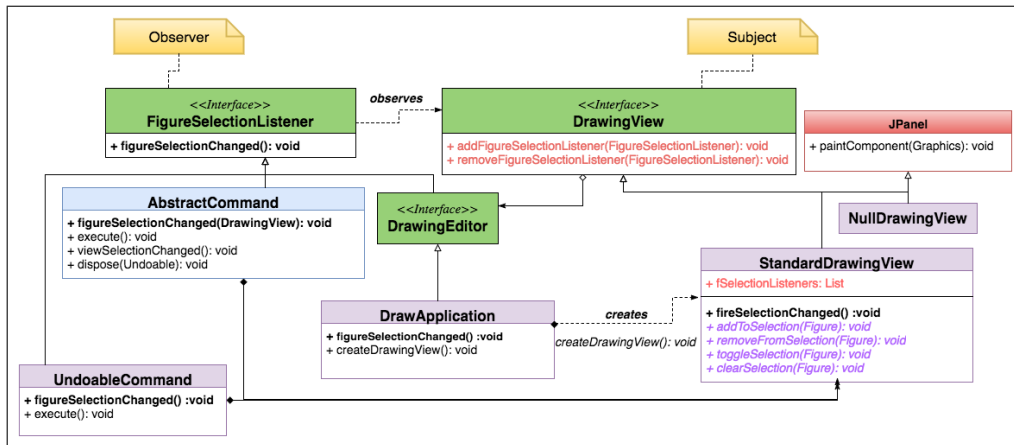

Figure 5.1: FigureSelectionListener in JHotDraw

As this figure illustrates, the `FigureSelectionListener` interface defines the *Observer* role. The classes that are interested in the changes of selection of figures in a `DrawingView` implement this interface. Accordingly, the `DrawingView` defines the *Subject* role, providing methods for adding and removing figure selection listeners. Practically, the only class that implements the *Subject* role is `StandardDrawingView`, while `NullDrawingView` has an empty implementation.

`StandardDrawingView` keeps the selection listeners in a list, the `fSelectionListeners`, and notifies them in the invocation of the `fireSelectionChanged` method. This method is called in the methods: `addToSelection`, `removeFromSelection`, `toggleSelection` and `clearSelection`, which indicate the change of figure selection. On the observers' side, the figure selection listeners implement the `figureSelectionChanged` method that is executed in case they have been notified by the subject.

Concluding, as described above, the "pattern code" of the observer pattern is scattered in many places, including the list of listeners on the subject, the add / remove methods, along with the pointcut methods that call the method which notifies the listeners.

5.5.2 Refactoring FigureSelectionListener in AJHotDraw

Marin et al. presented a refactoring of this concern in AJHotDraw [MMvD05a]. Their refactoring is illustrated in Figure 5.2.

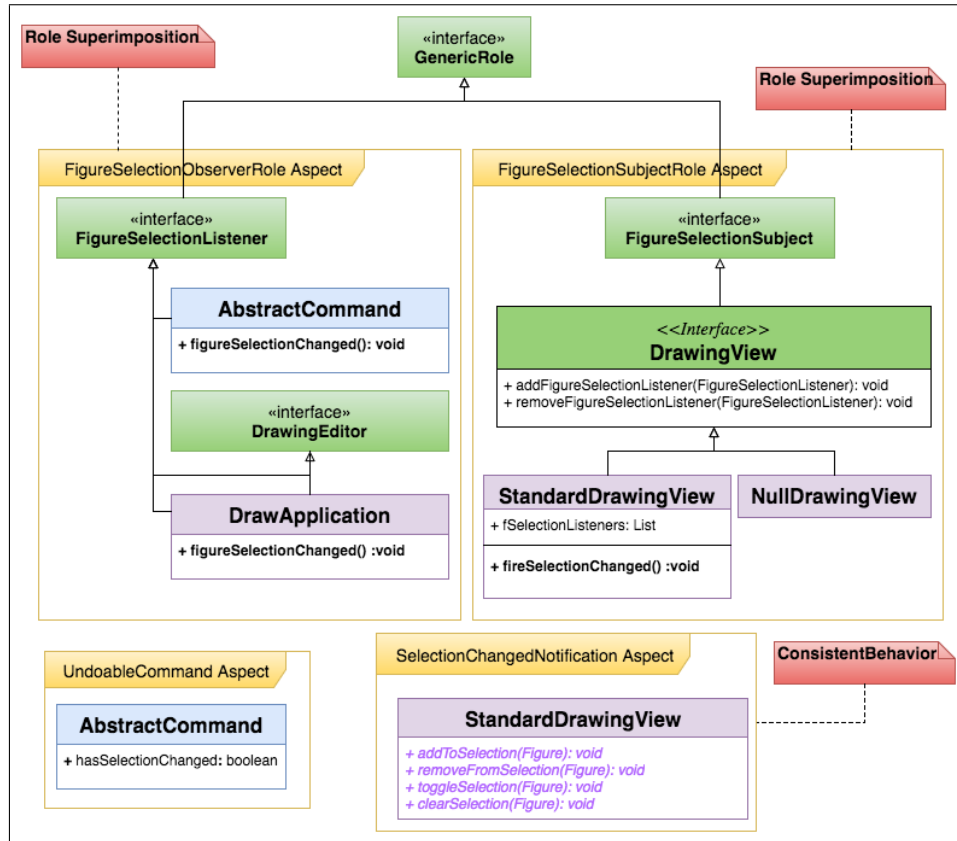Figure 5.2: FigureSelectionListener in AJHotDraw

In their proposed type-based refactoring, they have used two crosscut sorts, namely *role superimposition* and *consistent behavior*.

Role Superimposition

As defined by the authors [MMVD05b], "the role superimposition refers to the implementation of a specific secondary role or responsibility". In the case of `FigureSelectionListener`, they used it

twice, one for each of the roles. More specifically, they defined an abstract `GenericRole` and concrete roles, observer and subject which extend the abstract one.

Consistent Behavior

According to the authors[MMVD05b], "the consistent behavior sort implements a consistent behavior for a number of method elements that can be captured by a natural pointcut". In this case it is used to notify the *Observers* of the changes in the *Subject* object. More specifically, the methods `addToSelection`, `removeFromSelection`, `toggleSelection` and `clearSelection` are consistent behavior. They implement it as a pointcut in AspectJ shown in Listing 5.28.

```
public aspect SelectionChangedNotification {
  pointcut invalidateSelFigure(StandardDrawingView sdw) :
    (   withincode(boolean StandardDrawingView.addToSelectionImpl(Figure))
     || withincode(void StandardDrawingView.removeFromSelection(Figure)))
    && call(void Figure.invalidate())
    && this(sdw);

  pointcut clear_toggleSelection(StandardDrawingView sdw):
    (execution(void StandardDrawingView.clearSelection()) ||
     execution(void StandardDrawingView.toggleSelection(Figure)))
    && this(sdw);

  after(StandardDrawingView sdw): invalidateSelFigure(sdw) {
    sdw.fireSelectionChanged();
  }

  after(StandardDrawingView sdw): clear_toggleSelection(sdw) {
    sdw.fireSelectionChanged();
  }
}
```

Listing 5.28: AJHotDraw: Consistent Behavior in FigureSelectionListener

Benefits and Limitations

As the authors mention [MMvD05a], such refactoring allows the crosscutting elements to be addressed individually, which leads to a modular solution, and any deviations from the pattern implementation can be addressed separately.

However, the definition of pointcuts in order to capture the calls to the notifier is difficult when many consistent behavior instances occur. As Listing 5.29 shows, the original `clearSelection` method in JHotDraw calls `fireSelectionChanged` under specific conditions. Considering the AOP solution of AJHotDraw, Listing 5.28, this is not the case. In the pointcut definition, the pattern refactoring solution notifies the observers independently of the condition in the caller. Although, according to Marin et al. it is potentially harmless in this case, this implementation deviates from the behavior of the original JHotDraw, leading to a harmful, for the functionality, implementation. Finally, the problem of the unconditional call of a method in a pointcut is clearly a problem of the language. AspectJ mechanisms do not support such functionality. But can managed data solve this problem?

5.5.3 Refactoring FigureSelectionListener in ManagedDataJHotDraw

In JHotDraw's original code, the *observer* `DrawApplication` creates a new `StandardDrawingView` instance using the `createDrawingView` method. During the construction the `DrawApplication` passes itself to the constructor of the `StandardDrawingView` and this in turn adds it to its listeners list, using the `addFigureSelectionListener` method. This is shown in Line 7 of Listing C.3. Likewise, the rest

```
1  public void clearSelection() {
2    if (selectionCount() == 0) {
3      // avoid unnecessary selection changed event when nothing has to be cleared
4      return;
5    }
6    FigureEnumeration fe = selection();
7    while (fe.hasNextFigure()) {
8      fe.nextFigure().invalidate();
9    }
10   ...
11   fireSelectionChanged();
12 }
```

Listing 5.29: StandardDrawingView clearSelection Method

of the classes that implement the `FigureSelectionListener` interface, perform the same mechanism, adding or removing themselves from the `DrawingView` *Subject*. Consequently, the pattern code is scattered among its participants.

In this section we present our managed data refactoring of the `FigureSelectionListener` concern. Managed data implements aspects using data managers, by adding specifications to the data. For this case, we needed a similar mechanism to the *role superimposition* of AOP. This mechanism should be defined in a data manager that will produce managed data instances (managed objects) with a specific role. Additionally, the data manager has to support something similar to the *consistent behavior* for the pointcuts.

In detail, since the `DrawingView` is managed data, and it is the *Subject* to the *Listeners* of the *FigureSelectionListener* case, we can implement a data manager that attaches the *Subject* `MDStandardDrawingView`. Therefore, no *Subject* role specific code will be tangled with the `DrawingView`, but a data manager will attach this role later.

More specifically, we needed a data manager that performs the following:

1. Attaches the *Subject* role to the `MDStandardDrawingView` since this object implements the pattern. Initially, `MDStandardDrawingView` has no *Subject* role related code.

2. Enables the *Subject* to *add* and *remove* listener objects to and from itself.
   In this case the `FigureSelectionListener` instances.

3. Defines an *Action* that will be executed on the listeners in case of the *Subject*'s notification.
   In this case the `figureSelectionChanged` method.

4. Finally, it defines a pointcut for the consistent behavior that executes the actions on the listeners. In this case the `addToSelection`, `removeFromSelection`, `toggleSelection` and `clearSelection` methods.

Data manager

First, we can abstract the *Subject* role concern code in a separate data manager. As it is mentioned in Chapter 4, the role of a data manager class is to create a `MObject`, which interprets and handles a managed object instance. Therefore, we first need this `MObject`, namely `SubjectRoleMObject`, to implement our subject role specifications.

SubjectRole specifications

We define the functionality of the *SubjectRole* in an interface, shown in Listing 5.30. The subject role simply needs to add and remove listener objects to and from a managed object. Additionally, it defined a method `executeListenerActions` that allows the subjects to execute all the actions of their listeners.

```
1  public interface SubjectRole {
2    void addListener(Object listener, Action action);
3    void removeListener(Object listener);
4    void executeListenerActions();
5  }
```

Listing 5.30: SubjectRole Interface

Action

Additional to the listener object a *SubjectRole* has to define an *Action* for that listener. That *Action* determines the method which will be executed in that *Listener* in case a notification is retrieved from the *Subject*. As Listing 5.31 shows, this is simply a functional interface that represents an executable action.

```
1  @FunctionalInterface
2  public interface Action {
3    void execute();
4  }
```

Listing 5.31: Action Interface

SubjectRole MObject

Having the specifications of a Subject in place we need a data manager that implements them. This data manager defines a "role superimposition" of the subject role.

Role Superimposition

The implementation of the `SubjectRoleMObject` is presented in Listing 5.32. First, the SubjectRole-MObject extends the MObject, inheriting the functionalities of the base data manager, followed by the implementation of the `SubjectRole` specifications. By implementing the `SubjectRole` interface, the MObject has to implement the `executeListenerActions`, `addListener` and `removeListener` methods that have been provided by the subject role specifications. Having added a listener object along with its `Action` to be executed on each notification, the method `executeListenerActions` executes all the actions for each of the listeners.

```
1  public class SubjectRoleMObject extends MObject implements SubjectRole {
2    protected Map<Object, Action> listeners;
3    public SubjectRoleMObject(Klass schemaKlass, Object... initializers) {
4      super(schemaKlass, initializers);
5      listeners = new HashMap<>();
6    }
7    public void executeListenerActions() {
8      listeners.values().forEach(Action::execute);
9    }
10   public void addListener(Object listener, Action action) {
11     listeners.put(listener, action);
12   }
13   public void removeListener(Object listener) {
14     listeners.remove(listener);
15   }
16 }
```

Listing 5.32: SubjectRoleMObject

Consistent Behavior

Since we have implemented a form of *role superimposition* what is left is the *consistent behavior* pointcut. However, this functionality is application specific, therefore we can extend the abstract `SubjectRole` data manager with an application specific data manager that implements the consistent behavior. For practical reasons, we used an interface to define these specifications. As Listing 5.33 shows, a list of the methods that execute the `Action` for each listener are defined in the `FigureSelectionPointcut` interface.

```
public interface FigureSelectionPointcut {
  void addToSelection(Figure figure);
  void removeFromSelection(Figure figure);
  void toggleSelection(Figure figure);
  void clearSelection();
}
```

Listing 5.33: FigureSelectionPointcut Interface

FigureSelectionListenerSubjectRole MObject

Finally, having all of our specifications in place, we need to implement the actual MObject that uses them. Of course we need a new data manager that is application specific, namely FigureSelectionListenerSubjectRole. In particular, the FigureSelectionListenerSubjectRole MObject implements those specifications and provides its managed objects with the ability to use them. Additionally, this data manager extends the `SubjectRole` data manager, creating a stack of data managers. By stacking the data managers we have separated the application specific code, the consistent behavior in this case, from the more general subject role code.

Consistent Behavior Pointcut

Considering that an MObject is an Invocation Handler, every method invocation passes through that object first. By defining the pointcut in an interface and extending that interface in this MObject, we proxy the execution of the real object's methods, starting with MObject first. This allows the programmer to add functionalities in these methods which in other cases would scatter the real object's methods. Similarly to the AOP solution, the pointcut includes the three methods that call the `fireSelectionChanged` method. However, in managed data, we are not limited to a specific method of a class but to an `Action` passed for the specific listener. Invoking the `executeListenerActions` method on each of the methods defined by the pointcut, we have implemented the concern as a modular aspect. Listing 5.34 illustrates the code of the FigureSelectionListenerSubjectRole data manager.

Conditions in Pointcuts

As it has been seen from the AOP solution, during the pointcut definition, the language did not allow to add any kind of conditions or other functionalities based on the state of the object. However, since the MObject is proxied to a `MDStandardDrawingView` instance, the programmer can access the current state of the instance inside the data manager implementation. Therefore, the programmer can use the state of the program. In this case, the execution of the action on the listener is performed under a specific condition, (Line 27 of Listing 5.34), which is similar to the one defined on the original program 5.29.

```java
public class FigureSelectionListenerSubjectRoleMObject
  extends SubjectRoleMObject implements FigureSelectionPointcut {

  public FigureSelectionListenerSubjectRoleMObject(
    Klass schemaKlass, Object... initializers)
  {
    super(schemaKlass, initializers);
  }

  @Override
  public void addToSelection(Figure figure) {
    executeListenerActions();
  }

  @Override
  public void removeFromSelection(Figure figure) {
    executeListenerActions();
  }

  @Override
  public void toggleSelection(Figure figure) {
    executeListenerActions();
  }

  @Override
  public void clearSelection() {
    if (((MDStandardDrawingView) thisObject).selectionCount() > 0) {
      executeListenerActions();
    }
  }
}
```

Listing 5.34: FigureSelectionListenerSubjectRoleMObject
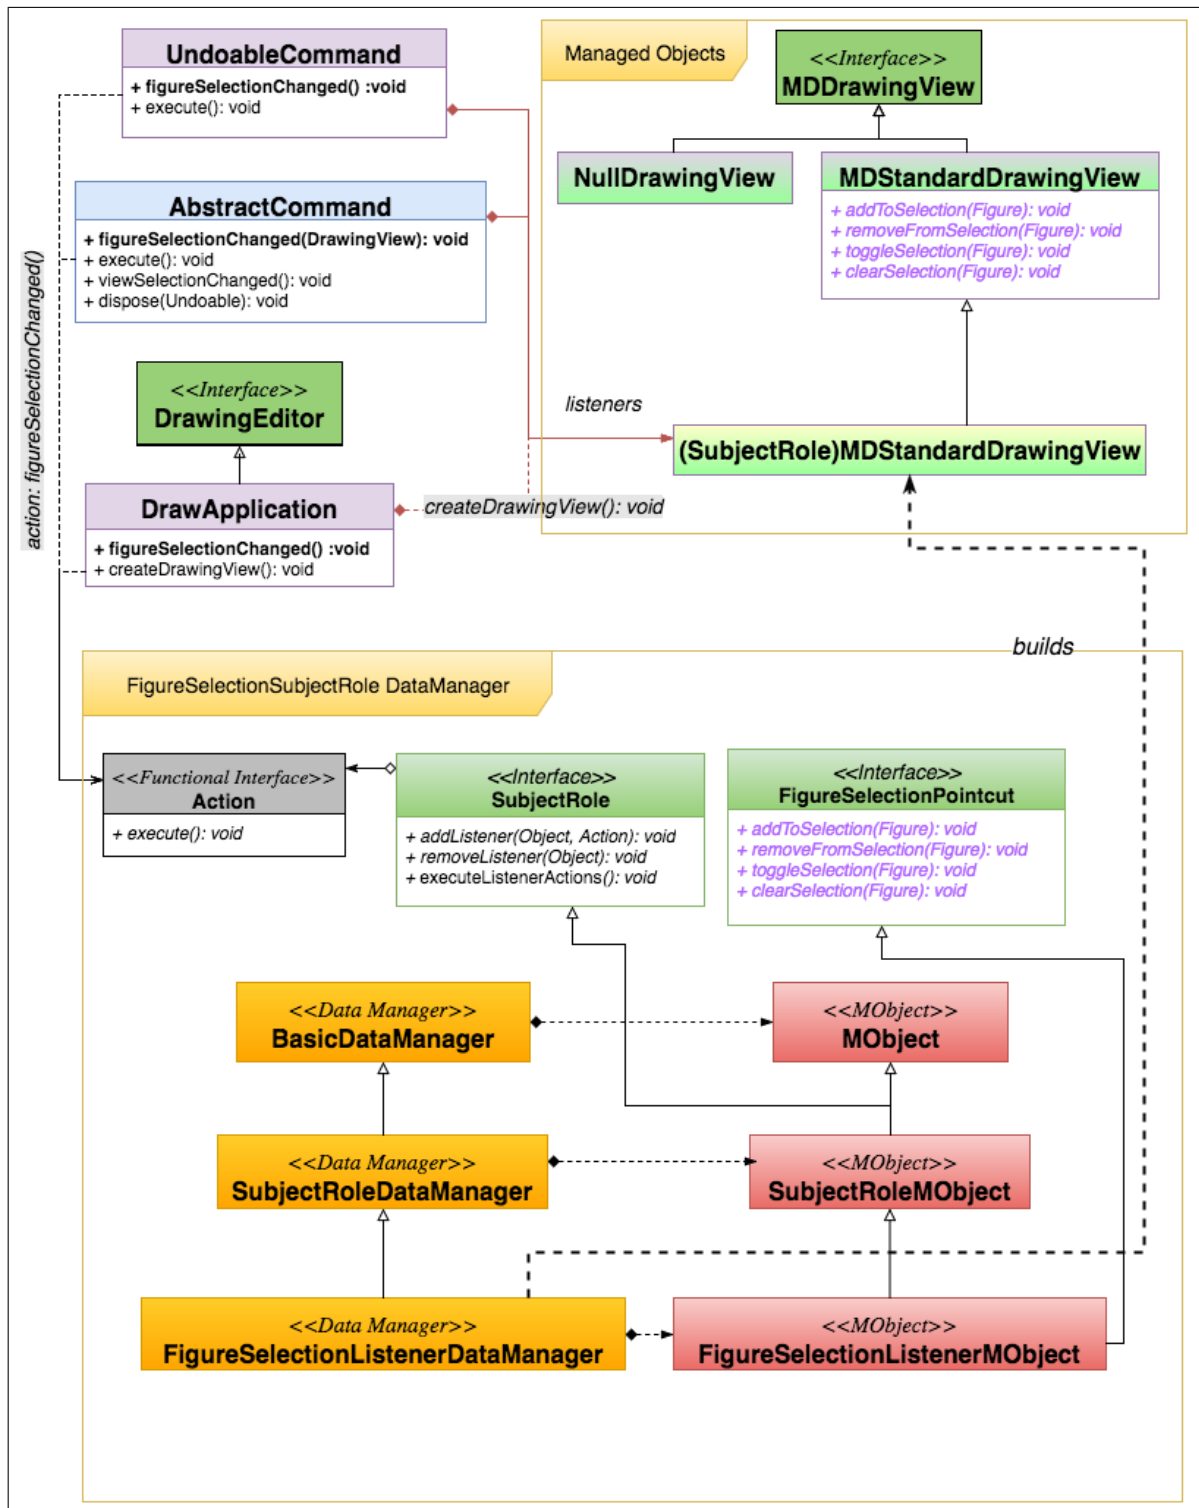
Figure 5.3: FigureSelectionListener in ManagedDataJHotDraw

5.5.4 Managed Data Refactoring

Figure 5.3 illustrates the refactored version of the FigureSelectionListener concern in ManagedData-JHotDraw. Comparing it to the original, Figure 5.1, it can be seen that, first, the list of listeners has been removed from the DrawingView. Next, the addListener and removeListener methods have

also been removed from the class. Every call of the `fireSelectionChanged` method in the pointcut methods has also been omitted. Finally, conditions on the pointcuts have been defined, something that is not supported by the AOP version, AJHotDraw.

The integration of the data manager was executed simply by using our schema factories and adding the listeners during construction. Most importantly, the behavior of the application remained equivalent to the original. ManagedDataJHotDraw conserved the behavior of JHotDraw, which we evaluated through its own test suite along with manual tests. Interestingly, by stacking data managers we manage to define aspects of our data in a modular way.

5.6 ChangeAttributeCommand: Undo Concern

The "Undo" functionality is used in several places in the original JHotDraw. Marin's fan-in analysis [MVDM04], identified about 30 undo activities defined for various elements of JHotDraw. For our assessment we focused on the refactoring of the undo concern in the *ChangeAttributeCommand* class. We choose the specific case since is the same that is used by Marin et al. on their undo refactoring in AJHotDraw [Mar04].

5.6.1 ChangeAttributeCommand in JHotDraw

The original Undo concern in *ChangeAttributeCommand* of JHotDraw is illustrated in Figure 5.4.
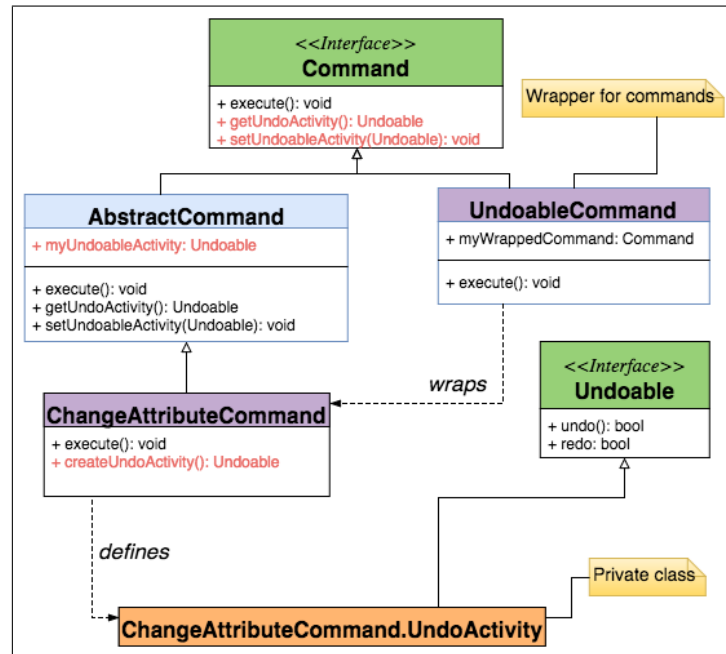
Figure 5.4: ChangeAttributeCommand in JHotDraw

A command of the *Command* pattern is represented by the `Command` interface in the figure. Some of the activities support the undo functionality, which is implemented in nested `UndoActivity` classes.

In the case of the *ChangeAttributeCommand*, the command is called when an attribute is applied to a figure. An attribute can be a color, a font, a url etc. When an attribute has been applied using a *ChangeAttributeCommand* object, the object defines its `UndoActivity` through the *UndoActivity* private class.

The role of the `UndoableCommand` wrapper class is to support repeated undo operations since it records the last executed commands of the wrapper class, in reverse order. In particular, this class acquires a reference to the undo activity associated with the wrapped command and it pushes it into a stack managed by an `UndoManager` [Mar04].

Therefore, as Marin et al. concluded [MVDM04] the "Undo" concern code is scattered in several places of the Command classes. First, the `myUndoableActivity` field in the `AbstractCommandClass` along with its accessors, `getUndoActivity` and `setUndoActivity`. Next, the private nested classes are implemented by the concrete commands that support undo. Moreover, the factory method, `createUndoActivity` creates instances of the private classes. Finally, the references to the before enumerated elements from non-undo related members.

5.6.2 Refactoring ChangeAttributeCommand in AJHotDraw

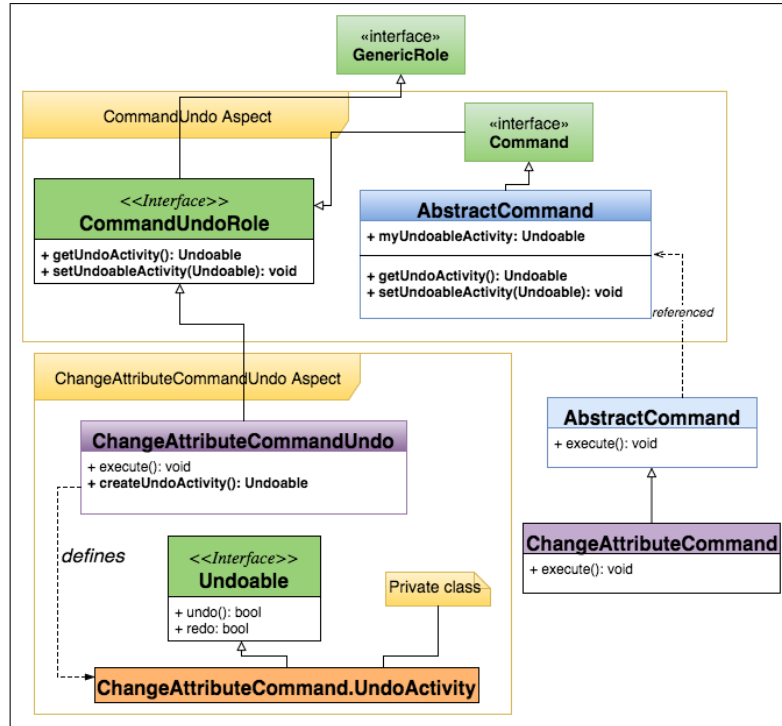The refactoring that Marin et al. proposed can be seen in Figure 5.5.

Figure 5.5: ChangeAttributeCommand in AJHotDraw

As the figure shows, a new aspect is created for the *ChangeAttributeCommand*. In this aspect the entire undo functionality is implemented and the undo code is removed from the actual *ChangeAttributeCommand* class. Additionally, the private class that implements the *UndoActivity* has moved to this aspect along with its factory method (*createUndoActivity*).

However, by convention, each aspect will consistently be named by appending "UndoActivity" to the name of its associated command class, to enforce the relation between the two. All the abstract undo functionality has been defined in a *CommandUndo* aspect. This aspect defines the undo as a role, the *Undoable* field and its accessors of the *AbstractCommand*.

Additionally, the change of the visibility of the methods introduced from the aspects is an issue. The visibility declared in the aspect refers to the aspect and not to the target class. However, the same problem occurs in managed data.

5.6.3 Refactoring ChangeAttributeCommand in ManagedDataJHotDraw

Finally, the aspect refactoring of the `ChangeAttributeCommand` undo concern in JHotDraw using managed data is presented in Figure 5.6.
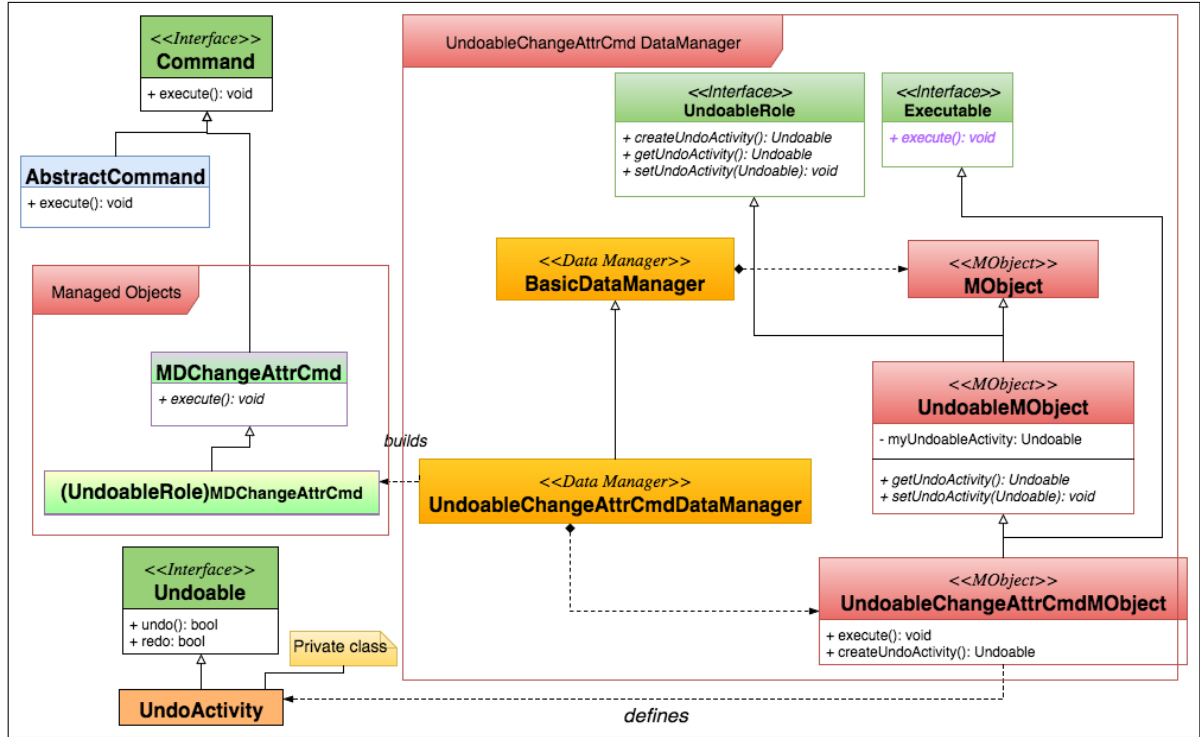
Figure 5.6: ChangeAttributeCommand in ManagedDataJHotDraw

To implement aspects using data managers in the `ChangeAttributeCommand` we first needed to migrate the original class to a managed data definition, namely *MDChangeAttrCmd*. The process was similar to the one described in Section 5.3.

In order to refactor the undo concern in managed data we needed a number of specifications that define the functionalities of our managed objects. As the AOP solution shows we needed something similar to *role superimposition*, but this time for the *Undoable* role. Additionally, we needed a way to define executables since our participant, in particular `ChangeAttributeCommand`, is a command.

In detail we needed a data manager that performs the following:

1. Attaches an *UndoableRole* to command instances. In this case the `MDChangeAttrCmd` instances.

2. Following the original design, this role should allow the objects to create their `UndoActivities` using a nested private class. *UndoableRole* should also provide accessors for its *UndoActivity*.

3. It defines an *Executable* specification, since every command has an `execute` method.

4. Finally, it defines a pointcut on the `execute` method since this is the place where the undo functionality for a specific command is implemented.

Data Manager

First, for abstraction's sake we have implemented the `UndoableRole` interface to an abstract class, namely `UndoableRoleMObject`. This data manager defines a `myUndoableActivity` field and implements its accessors given by the `UndoableRole` interface, including both `setUndoActivity` and `getUndoActivity`. However, no implementation for the `createUndoActivity` factory method is provided at this point. The reason is that this is a "command-specific" implementation and it has to be provided by the concrete classes. Since this is an abstract `MObject` no proxy factory is provided (data manager).

Next, we needed a `MObject` for the concrete command, in this case the `UndoableChangeAttrCmdMObject`. Since this has to inherit the `UndoableMObject`, it also has to implement the `createUndoActivity`

method. As mentioned, this is "command-specific" code, and following the original design, this functionality can be implemented in a nested class. As a result, an `UndoActivity` private nested class is defined in the `UndoableChangeAttrCmdMObject` class file. Instances of this class can be created using the `createUndoActivity` method.

Furthermore, the command is an executable object; therefore, this `MObject` has to implement the `Executable` interface. By doing so, the pointcut of the specific undo functionality can be implemented inside this `execute` method of the `UndoableChangeAttrCmdMObject`. A comparison of the `execute` method between the original and the refactored code can be seen in Listings 5.1 and 5.2.

```
1  public void execute() {
2    setUndoActivity(createUndoActivity());
3    getUndoActivity()
4      .setAffectedFigures(view().selection());
5    FigureEnumeration fe =
6        getUndoActivity().getAffectedFigures();
7    while (fe.hasNextFigure()) {
8      fe.nextFigure().setAttribute(fAttribute, fValue);
9    }
10   view().checkDamage();
11 }
```

Listing 5.1: Original execute method

```
1  default void execute() {
2    FigureEnumeration fe = view().selection();
3    while (fe.hasNextFigure()) {
4      fe.nextFigure().setAttribute(attribute(),value());
5    }
6    view().checkDamage();
7  }
```

Listing 5.2: Refactored execute method

As the two listings show, all of the *Undo* related functionalities of the `ChangeAttributeCommand` command have been removed from the `execute` method. This code is located at `UndoableChangeAttrCmdMObject` and more specifically at its execute pointcut. Listing 5.35 shows the pointcut code.

```
1  public void execute() {
2    setUndoActivity(createUndoActivity());
3    getUndoActivity().setAffectedFigures(thisObj.view().selection());
4  }
5  public Undoable createUndoActivity() {
6    return new UndoActivity(thisObj.view(), thisObj.attribute(), thisObj.value());
7  }
```

Listing 5.35: The execute Undo pointcut

5.6.4 Managed Data Refactoring

Comparing the original version with our refactored version several improved points can be observed. First, all of the *Undo* related functionalities of a command have been removed from the command's code and they have been externally attached by the *UndoRole* data manager. Next, following the original design, the creation of an `UndoActivity` instance is again defined inside a nested class; however, this time it is not in the command itself but it is inside a "command-specific" data manager. Finally, the `execute` method of a command is not aware of its undo functionality since it has been extracted to an undo dedicated pointcut inside a data manager.

5.7 Metrics

Concluding we have presented a new version of JHotDraw, the ManagedDataJHotDraw which refactors two CCC of the original system. We ensure that the new version preserved behavioral equivalence with the original, using the provided automated tests suite and manual tests. However, to present the relation between the original and our implementation we have collected a set of metrics [SGC+03].

5.7.1 Metrics Collection Details

First, those metrics measure the *Size*, which can show to what extend our version changed the original in terms of the size factor. Second, they measure *Cohesion* and *Coupling*, which shows how hard it is to understand, reuse and maintain the program. Finally, they measure *Separation of Concerns*, which presents the degree of scattering and tangling of the application . The *Separation of Concerns* metric is the most crucial for our research since this is the main criterion that we have focused on this refactoring.

For each concern refactoring on each system, we chose to collect data only for the participant classes so that our version would be comparable with the original. If all classes were included our results would be inflated strictly due to the higher number of classes. In addition, in the case of managed data, the abstract and reusable data manager classes have not been considered, i.e. the `SubjectRoleMObject` data manager. The reason is that the data managers are reusable code and independent from the application's logic. Certainly, any (low level) data managers that are application-specific, like the point-cut data managers, are measured, i.e. the `FigureSelectionListenerMObject`.

The metrics WOC, CBC, LCOO and DIT are similar to Chidamber and Kemerer's metrics and easy to be collected using simple software analysis tools. Likewise, these tools can easily measure LOC, in our case we have used the Non-Comment Lines Of Code (NCLOC) for reliability.

The NOA metric counts the attributes of a class. The number of attributes in both OOP and AOP systems it is simply the number of fields. However, NOA in the managed data system has a significant difference. In managed data schema definitions everything is method declarations, even the fields; therefore, in order to count NOA we need to count the method declarations which are not `default` and therefore represent field declarations.

Finally, note that no data have been collected for the AJHotDraw version but only for the original JHotDraw and the managed data version, MDJHotDraw. The reason is that, although those metrics support aspect abstractions in their definition [SGC+03], the comparison is not well-grounded. Establishing reliable comparison between programs written in Java and programs written in AspectJ (and Java), could be misleading. Because those languages differ in their abstraction mechanisms. JHotDraw and MDJHotDraw are both written in pure Java; using the same abstractions. However, in managed data the design has been changed significantly and thus, we compare the two systems in order to show how this design refactoring affects the original implementation.

5.7.2 Data Collection

For the data collection we used the `MetricsReloaded` [4]plug-in for IntelliJ IDE[5]. Using this tool we managed to collect the following metrics[6]: LOC, WOC, CBC, LCOO and DIT. Table 5.1 presents the metrics collected for the systems under evaluation.

	Metrics	ObserverPattern (FigureSelectionListener)		Undo (ChangeAttributeCommand)	
		JHotDraw	**MDJHotDraw**	**JHotDraw**	**MDJHotDraw**
	CDC	8	4	5	3
SoC	**CDO**	21	8	16	8
	CDLOC	46	6	32	12
Coupling	**CBC**	170	184	128	128
	DIT	18	16	10	11
Cohesion	**LCOO**	75	64	14	12
	VS	8	10	5	5
Size	**LOC**	1793	1878	1052	1063
	NOA	53	53	29	27
	WOC	383	368	198	199

Table 5.1: Metrics

As presented in the table, the *coupling* metric of the two systems is very similar. In the case of

---

MDJHotDraw the number has increased since the data managers are coupled with their specifications but not with the application, while in the original JHotDraw the crosscutting concerns are coupled with the application components. The *size* of the two systems is also similar. For MDJHotDraw it is slightly bigger because of the migration to managed data. The *cohesion* of MDJHotDraw is moderately smaller since during the managed data migration we have tried to improve cohesion explicitly. Finally, the *separation of concerns* of MDJHotDraw is notably higher. This is because the concern code is not scattered and tangled like in the JHotDraw case, but it has moved to data managers. The detailed results are presented in Appendix D and discussed in the following chapter.

Chapter 6

Evaluation

Having presented evidence that our framework is able to solve the problem of CCC (Chapters 3 and 5) we will now evaluate our claims. In this chapter we present an evaluation of our contributions in implementing aspects with managed data, in relation to our research questions.

6.1 Research Questions and Answers

In Chapter 1 we stated three research questions which were the main focus of this thesis. These are answered as follows:

How to implement managed data in a static language? From our managed data implementation in Java, presented in Chapter 4, we conclude that it is possible to implement managed data in a static programming language by using Reflection and Dynamic Proxies.

Can managed data solve the problem of crosscutting concerns? We argue that by using data managers, managed data can implement aspects of data in a modular way. As presented in both Chapters 3 and 5, managed data can handle aspects and solve the problem of Cross Cutting Concerns. Additionally, we have collected a set of metrics that assess our implementation in relation to the original. In this chapter we discuss the results of those metrics extensively.

To what extent can managed data handle an inventory of aspects in the JHotDraw framework, compared to the original and the Aspect Oriented implementation? The results of our aspect refactoring in Chapter 5 show that managed data can handle aspects of JHotDraw. In addition, we claim that our solution extends some of the capabilities of the AJHotDraw implementation. However, in this chapter, we provide further evaluation of our aspect refactoring and we compare it with the Aspect Oriented Programming solution in terms of a list of "modularity properties" and metrics.

6.2 Assessment Framework

In this section we discuss the summary of those results presented in Table 5.1.

6.2.1 Metrics Interpretation

Coupling. Although the numbers of Depth of Inheritance Tree are similar in the two systems, the reasons are not the same. In the case of the JHotDraw the classes have higher DIT because they implement interfaces that define concerns. This is not the case in MDJHotDraw; however, the high number of DIT is due to the data managers stacking process. Additionally, the specifications are described in interfaces implemented by data managers which also increases the DIT. However, the coupling is between specifications and data managers and not CCC and components.

The Coupling Between Components metric is also similar in both versions. Again, the reasons differ for each case. CBC in JHotDraw increased from the coupling between main components and CCC. In the case of MDJHotDraw there is not such coupling but there is coupling between data managers and their specifications.

The coupling overall number in MDJHotDraw is similar to JHotDraw; however, for totally different reasons. JHotDraw components are coupled with the CCC, while MDJHotDraw data managers are coupled with their specifications. Therefore, data managers are not coupled with the application logic since they are abstract and reusable. This solved the evolution paradox [TBG03] problem observed in AOP, where the aspects are closely coupled with the components since the aspects have to know about the components.

Size. The Vocabulary Size metric is higher in the case of MDJHotDraw since, although the concern-related components have been removed, more components have been added for the managed data migration. Those components are schema factories and helpers that altogether migrate the system to managed data.

As a result, the Lines Of Code metric is higher for the MDJHotDraw version since more components have been added.

The Weighted Operations per Component metric is similar in both cases. The reasons are that no significant changes have been made in terms of component complexity and the code remained similar to the original one. Small fluctuations on the numbers indicate that complexity increased or decreased by the data managers code in the MDJHotDraw case.

Finally, the Number of Attributes metric is almost identical in both cases. The reason is that attributes that were related to the concern in JHotDraw have now moved to specifications and fields in low-level data managers in MDJHotDraw.

In general, the size of the two systems is similar with a small increase in the MDJHotDraw owing to the fact that a few components have been added during the migration to managed data process.

Cohesion. The Lack of Cohesion in Methods metric values are lower on the MDJHotDraw case. Although the implementation of the point-cuts in managed data is performed using interfaces and this affects the LCOM metric of the concrete implementation classes, the value is still lower than JHotDraw. This is mainly because of the managed data migration. During this process we tried to improve the cohesion of the managed data definition by separating non-managed data from managed data as much as possible.

Separation Of Concerns. Finally, the following metrics measure the scattering and tangling of code in the two versions. First, both Concern Diffusion over Components and Concern Diffusion over Operations metrics measure the code scattering in the systems. The differences in these two metrics are significant since the managed data version explicitly focuses on removing scattered code from the original and moving the concern code in data managers. Especially the CDO metric, which measures the operations that assist the implementation of a CCC, is notably lower because those operations are now located in data managers. Second, the Concern Diffusion over LOC metric measures the code tangling of the systems. The MDJHotDraw has a remarkably lower number owing to the fact that every concern-related code has been removed completely from the application classes. The only place that CDLOC exists is in the integration points of the clients classes. More specifically, the places that define the usage of this concern by using the proper data manager.

In general, scattering and tangling metrics are significantly low in the case of MDJHotDraw. This is because we have explicitly moved the concern-code from the classes to reusable data managers.

6.3 Modularity Properties

In order to make an assessment of our research in comparison with the AOP version, we applied the same software quality metrics presented by Hannemann et al. [HK02] in their design patterns refactoring. More specifically, the authors refer to those metrics as "closely-related modularity properties" and use them to analyze and evaluate their design pattern solutions in AspectJ.

Since our solution of the observer pattern concern in JHotDraw refers to the same issue that Hannemann et al. [HMK05] describes, we concluded that using the same modularity metrics will lead to a more reliable assessment of our solution and a more definitive assessment of our designs. The focal modularity properties are: *Locality*, *Reusability*, *Composition transparency* and *(Un)pluggability*.

6.3.1 Modularity Properties in the Observer Pattern

In their *Observer* pattern refactoring in AspectJ, Hannemann et al. [HMK05], along with Marin et al. [MMvD05a] who used the same refactoring method in AJHotDraw, assess their results using these four modularity properties.

Locality

As shown in the AOP solution, the code that implements the observer pattern is in concrete observer aspects, nothing is scattered in the participant classes. Thus, all the participants are free of the pattern context and there is no coupling between the participants [HMK05]. However, as the AOP evolution paradox suggests [TBG03], the aspects themselves are strongly coupled to the application's code.

Likewise, in the *managed data* version, *Subject* and *Observer* classes are pattern agnostic. However, those classes are not aware of their pattern properties since they have been attached later by a subject role data manager.

Reusability

In the AOP version, the core pattern code is abstract and reusable. They have implemented abstract aspect code that explicitly refers to an observer pattern generalized implementation. This has been defined as an abstract aspect and it can be reused and shared across multiple observer pattern instances [HMK05]. That leads to a modular solution, although the concrete aspects still are highly coupled with the application's code.

Similarly, in the *managed data* version, we have implemented the observer pattern as a reusable data manager. More specifically, the `SubjectRole` data manager is an abstract observer pattern that can be reused in any case. Moreover, the definition of this data manager is independent from the object's structure. No coupling between the application and the observer pattern implementation exists. The only "case-specific" functionality is the consistent behavior, which is similar to the one presented in the AOP solution, although, our solution supports conditional pointcuts. Additionally, the managed data version defines the pointcut in an explicit interface, separating the general observer from the specific consistent behavior.

(Un)pluggability

In the AOP version, since *Subjects* and *Observers* do not need to be aware of their role in any pattern instance, it is possible to switch between using a pattern and not using it in the system [HMK05].

The *managed data* version, takes this a step further. It allows the programmer to "plug" or "unplug" specifications of the pattern by simply creating managed objects, using the preferable data manager with or without the specifications for the pattern. One can switch to the pattern implementation by creating objects provided by the subject role data manager, or simply by the basic data manager, which would lead to a non-observable object. Furthermore, this (un)pluggability can be enabled in data managers composition level, by stacking data managers, something we have applied in the subject role and pointcut case.

Composition transparency

In the AOP version, a pattern's participants are not coupled to the pattern, neither to the abstract aspects. Therefore, if a *Subject* or an *Observer* takes part in multiple observing relationships, their code does not become more complicated [HMK05]. However, the composition becomes more complicated in the join-points, since such an aspect implementation has to include multiple roles.

In the *managed data* version, we can implement such a thing in various ways. First, we can choose the AOP approach, by implementing a generalized data manager for the pattern and concrete data managers for each implementation. Alternatively, we could implement a data manager which implements a set of specifications that result to the composition of patterns. Such solution would lead to a separate specification definition that describes a composition of patterns. Consequently, the final system will be simple and (un)pluggable since the composition is transparent in a data manager.

6.3.2 Unpluggability of the Undo Concern

In order to evaluate the CCC refactoring of "Undo" Marin [Mar04] used the (Un)pluggability property. The author groups the complexity of the commands based on two criteria. First, on the degree of *tangling* of the undo setup in the command's logic, specifically the activity's `execute()` method. Second, on the impact of removing the undo-related part from its original site, which can be estimated by the number of references to the factory method and to the methods of the nested undo activity. Thus, the (un)pluggability property gives a measure of how clearly the concern is distinguished in the original code and is a good estimate of the refactoring costs.

In the AOP version of the `ChangeAttributeCommand` the refactoring is considered successful. First, the `UndoActivity` nested class, accompanied by its factory method, is removed from the command's code. Next, the accessors of the `UndoActivity` are inherited from top level classes and not overridden locally. Finally, the undo related code in the nested classes is *unpluggable* and suitable for extraction and refactoring.

Likewise, the managed data version followed the same unpluggability principles as the AOP version. Following the similar design we managed to achieve the same level of unpluggability. However, in this case the methods related to the crosscutting functionality of undo have not been defined by top level classes but by a set of stacked data mangers. Additionally, in managed data the `UndoActivity` is not defined based on naming conventions like in AspectJ, but through a nested private class definition dedicated to the command-specific data manager.

The overall assessment of the modularity properties presented in this section is summarized in Table 6.1.

Modularity Property	AJHotDraw	MDJHotDraw
Locality	+ CCC implemented in aspects	+ CCC implemented in data managers
Reusability	+ Abstract aspects - Pointcuts for concern-specific - Not Flexible (Name conventions, no conditions)	+ Abstract data managers - Pointcuts for concern-specific + Flexible (No conventions, state in DMs)
(Un)pluggability	+ Easy to plug and unplug (Defined in aspect)	+ Easy to plug and unplug (Change data manager)
Composition Transparency	+ Easy to be part on more than one patterns - Complicated in join-points (multiple role aspects)	+ Easy to be part on more than one patterns + Easy on join-points (stack data managers)

Table 6.1: Modularity Properties

6.4 Discussion

In this section we provide a discussion of the overall findings during the development of this thesis.

6.4.1 Managed Data Implementation

One of the advantages of our managed data framework is the simplicity of its usage. The programmer simply needs to define the schemas, followed by their factories and data managers, can easily write a program using them. The framework takes care of the *parsing*, *interpretation*, *type checking* and any other underline mechanisms. Moreover, it uses Java concepts, which makes it safer in terms of type checking and definitions resulting to an easier, for Java developers, adaptation. Furthermore, by being a self-describing language it is no longer bounded to the Java constructs transforming everything into managed data. Finally, the effortless mechanism of stacking data managers makes it significantly modular on every level, meta or not.

However, there is a number of implementation issues described in Section 4.5.

Another issue that arises is that integration in existing systems is complicated considering every model has to be redefined as a schema and every functionality has to be reimplemented in data managers. An existing system integration is presented in Chapter 5.

6.4.2 Managed Data Evaluation

Managed data is able to successfully handle aspects in a system. However, the following issues should be acknowledged:

Modularity

First, the modularity of data managers and their abstract way of data manipulation makes the implementation of aspects very simple and extensible. More specifically, implementing *components* (non-crosscutting concerns) as managed data and *aspects* (crosscutting concerns) by using specific data managers, one completely separates the two concepts. Contrary to the AOP version, managed data does not couple the aspects with the application's components code. The aspects are reusable data managers that can be plugged / unplugged and used in various concerns. Certainly, there are application specific parts, such as pointcut definitions; however, by *stacking* data managers we can support the application specific pointcuts at the lowest level. As a result, we can claim that managed data does not entirely solve the **evolution paradox** [TBG03], but it exceeds AOP system in modularity.

Flexibility

Managed data can be used similarly to AOP; however, the flexibility of the language takes it a step further in some cases. The design of the application is not bounded by the language, like in AspectJ's case.

AspectJ has limitations that can be one overcome by using managed data. In particular, as shown, data managers allow to access the current object's state, while AspectJ's pointcuts implementation do not. An object's state can be used for implementing functionalities depending on the current state of an object, for instance conditions in pointcuts.

Furthermore, AOP tries to discriminate methods based on some common structural properties such as particular coding conventions. For instance, AspectJ's mechanisms do not allow introduction of nested classes; therefore in the case of the undo refactoring, the post-refactoring association was only an indirect one, based on naming conventions ("`UndoActivity`"). This is a weaker connection than the one provided by the original solution. Additionally, what happens in case the conventions followed do not agree with the rules, or in case developers do not follow the desired conventions? In other words, it is very difficult to capture the required join-points for the aspect weaver in a general and extensible way. Managed data on the other hand do not require such conventions. Since we have defined our data, its structure is accessible via the `schemaKlass` of the `MObject`, therefore we do not need any kind of conventions to determine the structure of properties. Moreover, private classes in managed data is as easy to be defined as in the original version since `MObject` is just a class. Therefore, no code conventions are needed.

Performance

In practice, since managed data has been implemented using Java reflection and Dynamic Proxies, it is unfavorable for applications that need performance to use it. AOP, and specifically the *aspect weaving* process, provides an oblivious way of dealing with aspects. The weaver produces static Java code, which is then compiled in JVM and can be optimized. On the contrary, data managers dynamically analyze the schemas through reflection, which makes it a lot harder for the compiler to optimize. More specifically, even though the HotSpot JVM has one of the best just-in-time compilers, Java's dynamic proxies introduce 6.5x overhead [MSD15]. Thus, AOP performance is much higher than managed data.

Migration and Integration

Finally, both the migration and the integration of an existing application in the two cases has some trade-offs.

On one hand, the integration of an AOP version requires a whole new language (e.g. AspectJ) in line with the new concepts of the AOP paradigm. A developer has to implement *aspects*, *advices*, *join-points* and *pointcuts*. In addition, the programmer has to setup AspectJ to the IDE of preference along with the programming environment. After the environment has been set up, the migration of an existing application to the AOP paradigm is relatively easy. Finally, a developer has to migrate the application using the aspect oriented language's concepts and the weaver will do the job.

On the other hand, the integration of managed data is very simple since it is pure Java. No new language is required, no any additional IDE setup nor programming environment. Additionally, although managed data introduces new concepts, it still uses the Java syntax, e.g. interfaces and annotations; therefore, the integration is a lot simpler. It is relatively easy for a developer, who is familiar with Java, to learn the managed data concepts. However, the migration of an existing application to managed data can be time consuming since a programmer has to migrate it from normal Java class definition to managed data schemas definition (interfaces). Additionally, some limitations arise from this migration because interfaces and classes do not explicitly support the same functionalities. As seen, the framework has some limitations such as Java keywords, encapsulation of the methods in interfaces or class inheritance. However, such definitions could be defined using annotations in future work.

6.5   Threats to Validity

In this section we present a set of validity criteria including *Construct*, *Internal*, *External* validity and *Reliability* [ESSD08]. First, regarding the construct validity of this research, one should keep in mind that our results are compared to another research's findings. The AJHotDraw implementation may not be indicative of the overall AspectJ's capabilities in aspect refactoring nor was altered by us. In other words our results focused on the comparison of our aspect refactoring in managed data with Marin's et al. AJHotDraw aspect refactoring.

The internal and external validity of this thesis is satisfactory since our focus was narrowed to a representative case and predefined and tested metrics. However, the AJHotDraw provides a more complete aspect refactoring of the original version. For instance, in the undo concern case, AJHotDraw refactors all of its instances inside JHotDraw. In our case we only refactored a part of the undo concern, in order to show how our framework will handle this aspect; a limited approach when compared to AJHotDraw's holistic implementation. We should keep in mind that the metrics we used have been introduced for both OOP and AOP systems. However, in this thesis we used them in managed data implementation. Furthermore, in order to evaluate the behavioral conservation of our refactoring we used JHotDraw's test suite. Even though our code passes the tests, the suite is not complete and breaks in some files, even in the original version. We proposed the creation of a new test suite, like AJHotDraw did with the TestJHotDraw project, as a potential solution to this problem.

Additionally, one can argue that JHotDraw's CCC could be refactored simply by using dynamic proxies, without the whole Managed Data modeling structure. This is something that can be researched more; however, managed data besides the first class modeling they provide, they also have a modular design of aspect implementation. This design makes it easy to write modular aspects of data using dynamic proxies.

Finally, the reliability of this research can be problematic. More specifically, the refactoring is strictly dependent on the programmer's design and unless our design is used when replicating this research, a different assessment could be reached.

Chapter 7

Conclusion

In this research we have presented a Managed Data implementation in Java using its reflection API and dynamic proxies. By doing so, we have implemented a powerful approach to data abstraction that is otherwise hard-coded in the programming language. Our first contribution is the creation of a framework, the JavaMD, which is an open-source project of managed data implementation in Java.

Having the managed data implementation in place, we refactored an existing Object Oriented use case, the JHotDraw. Although this use case is considered as a well-designed OOP system, it faces the problem of the CCC. Thus, this system was refactored using managed data in order to solve the CCC problem. We migrated some main components of JHotDraw to managed data, removed the observer pattern and undo CCC and, finally, used data managers to implement them. This refactoring led to a new version of JHotDraw, the ManagedDataJHotDraw, which solves the problem of some main CCC of the original application. Our second contribution is the ManagedDataJHotDraw framework as an open-source project.

During the assessment of our refactoring we collected a number of metrics that we used in order to evaluate our refactoring in regard to the original application. Moreover, we extensively presented the refactoring process of JHotDraw and compared with AJHotDraw, the Aspect Oriented implementation of JHotDraw. Finally, by presenting a set of metrics and a number of modularity properties, we assessed our results in relation to the original and the Aspect Oriented version. This extensive evaluation of our refactoring leads to our third contribution.

Overall, we showed that managed data can be implemented in a static language and it can handle aspects by using data managers for concern implementation. Moreover, it extends some capabilities of AspectJ and deviates from the problem of coupling between aspect definition and components, leading to a modular and flexible way of aspect implementation. Although there are several challenges to be solved, such as the significant performance overhead of reflection, our work shows how managed data gives the programmer control over the fundamental mechanisms for creation and manipulation of data. In addition, it demonstrates how managed data can be used to solve the problem of CCC in existing or new systems.

Acknowledgments

It all started when my supervisor, Tijs van der Storm, presented us the classic Aspect Oriented Programming paper during our software construction course of this Master's program. I was always thinking that scattering was annoying while tangling looked very bad, but I did not know the name. The Aspect Oriented Programming paper explained it to me explicitly. First thing I did was to go to Tijs and ask him if AOP really works, since I was not really convinced. How can no one use this technique if it really works? Tijs then explained to me that he and William R. Cook worked together on something different, that would solve that problem in a different way, the Managed Data. That day I left his office with a list of papers discussing AOP and Managed Data. I started implementing it in Java the same day in order to see if it actually works. Well, it worked, and I got excited. This was the time I knew I wanted to do my research on this.

For this thesis, I would first like to express my sincere gratitude to Tijs van der Storm, my supervisor, who introduced me to this idea. He helped me with the problems I faced during this process. I will never forget the self-describing schema process; Tijs had the patience to explain it to me a number of times. I also want to thank Tijs because he was a real inspiration to me during my master's program, in particular his way of thinking, his knowledge and of course his perfectionism. He changed my way of programming and thinking about programming entirely.

Moreover, I would like to thank my colleagues and friends from CWI for the interesting and inspiring conversations. However, my biggest support was from Ifigeneia. The writing of this thesis would never be complete if she was not on my side. I am very grateful to Ifigeneia for her precious constant help, irreplaceable support and her interest on the subject during its completion.

Bibliography

[AD08]   Michael D Adams and R Kent Dybvig. Efficient nondestructive equality checking for trees and graphs. In *ACM Sigplan Notices*, volume 43, pages 179–188. ACM, 2008.

[BVDVET05]   Magiel Bruntink, Arie Van Deursen, Remco Van Engelen, and Tom Tourwe. On the use of clone detection for identifying crosscutting concern code. *Software Engineering, IEEE Transactions on*, 31(10):804–818, 2005.

[CK94]   Shyam R Chidamber and Chris F Kemerer. A metrics suite for object oriented design. *Software Engineering, IEEE Transactions on*, 20(6):476–493, 1994.

[ESSD08]   Steve Easterbrook, Janice Singer, Margaret-Anne Storey, and Daniela Damian. Selecting empirical methods for software engineering research. In *Guide to advanced empirical software engineering*, pages 285–311. Springer, 2008.

[Eug06]   Patrick Eugster. Uniform proxies for java. *ACM SIGPLAN Notices*, 41(10):139–152, 2006.

[FAG+08]   Eduardo Figueiredo, Claudio Sant Anna, Alessandro Garcia, Thiago T Bartolomei, Walter Cazzola, and Alessandro Marchetto. On the maintainability of aspect-oriented software: A concern-oriented measurement framework. In *Software Maintenance and Reengineering, 2008. CSMR 2008. 12th European Conference on*, pages 183–192. IEEE, 2008.

[FFI04]   Ira R Forman, Nate Forman, and John Vlissides Ibm. Java reflection in action. 2004.

[Fow09]   Martin Fowler. *Refactoring: improving the design of existing code*. Pearson Education India, 2009.

[Gam95]   Erich Gamma. *Design patterns: elements of reusable object-oriented software*. Pearson Education India, 1995.

[GSC+03]   Alessandro Garcia, Cláudio SantAnna, Christina Chavez, Viviane Silva, Carlos Lucena, and Arndt von Staa. Agents and objects: An empirical study on the design and implementation of multi-agent systems. In *Proc. of the SELMAS*, volume 3, pages 11–22, 2003.

[HK02]   Jan Hannemann and Gregor Kiczales. Design pattern implementation in java and aspectj. In *ACM Sigplan Notices*, volume 37, pages 161–173. ACM, 2002.

[HMK05]   Jan Hannemann, Gail C Murphy, and Gregor Kiczales. Role-based refactoring of crosscutting concerns. In *Proceedings of the 4th international conference on Aspect-oriented software development*, pages 135–146. ACM, 2005.

[KBJV06]   Ivan Kurtev, Jean Bézivin, Frédéric Jouault, and Patrick Valduriez. Model-based dsl frameworks. In *Companion to the 21st ACM SIGPLAN symposium on Object-oriented programming systems, languages, and applications*, pages 602–616. ACM, 2006.

[KDRB91]   Gregor Kiczales, Jim Des Rivieres, and Daniel Gureasko Bobrow. *The art of the metaobject protocol*. MIT press, 1991.

[KHH+01]   Gregor Kiczales, Erik Hilsdale, Jim Hugunin, Mik Kersten, Jeffrey Palm, and William G Griswold. An overview of aspectj. In *ECOOP 2001Object-Oriented Programming*, pages 327–354. Springer, 2001.

[KLM+97]   Gregor Kiczales, John Lamping, Anurag Mendhekar, Chris Maeda, Cristina Lopes, Jean-Marc Loingtier, and John Irwin. Aspect-oriented programming. In *ECOOP'97Object-oriented programming*, pages 220–242. Springer, 1997.

[LL00]   Martin Lippert and Cristina Videira Lopes. A study on exception detection and handling using aspect-oriented programming. In *Proceedings of the 22nd international conference on Software engineering*, pages 418–427. ACM, 2000.

[LvdSC12]   Alex Loh, Tijs van der Storm, and William R Cook. Managed data: modular strategies for data abstraction. In *Proceedings of the ACM international symposium on New ideas, new paradigms, and reflections on programming and software*, pages 179–194. ACM, 2012.

[Mar04]   Marius Marin. Refactoring jhotdraws undo concern to aspectj. In *Proceedings of the 1st Workshop on Aspect Reverse Engineering (WARE 2004)*, 2004.

[MM]   Marius Marin and Leon Moonen. Ajhotdraw: A showcase for refactoring to aspects.

[MMvD05a]   Marius Marin, Leon Moonen, and Arie van Deursen. An approach to aspect refactoring based on crosscutting concern types. In *ACM SIGSOFT Software Engineering Notes*, volume 30, pages 1–5. ACM, 2005.

[MMVD05b]   Marius Marin, Leon Moonen, and Arie Van Deursen. A classification of crosscutting concerns. In *Software Maintenance, 2005. ICSM'05. Proceedings of the 21st IEEE International Conference on*, pages 673–676. IEEE, 2005.

[MSD15]   Stefan Marr, Chris Seaton, and Stéphane Ducasse. Zero-overhead metaprogramming: Reflection and metaobject protocols fast and without compromises. In *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation*, pages 545–554. ACM, 2015.

[MVDM04]   Marius Marin, Arie Van Deursen, and Leon Moonen. Identifying aspects using fan-in analysis. In *Reverse Engineering, 2004. Proceedings. 11th Working Conference on*, pages 132–141. IEEE, 2004.

[Par72]   David Lorge Parnas. On the criteria to be used in decomposing systems into modules. *Communications of the ACM*, 15(12):1053–1058, 1972.

[PSH04]   Polyvios Pratikakis, Jaime Spacco, and Michael Hicks. Transparent proxies for java futures. In *ACM SIGPLAN Notices*, volume 39, pages 206–223. ACM, 2004.

[RM02]   Martin P Robillard and Gail C Murphy. Concern graphs: finding and describing concerns using structural program dependencies. In *Proceedings of the 24th international conference on Software engineering*, pages 406–416. ACM, 2002.

[SCL+12]   T Storm, WR Cook, A Loh, K Czarnecki, and G Hedin. Object grammars: Compositional & bidirectional mapping between text and graphs. 2012.

[SGC+03]   Cláudio SantAnna, Alessandro Garcia, Christina Chavez, Carlos Lucena, and Arndt Von Staa. On the reuse and maintenance of aspect-oriented software: An assessment framework. In *Proceedings of Brazilian symposium on software engineering*, pages 19–34, 2003.

[SGP04]   David C Shepherd, Emily Gibson, and Lori L Pollock. Design and evaluation of an automated aspect mining tool. In *Software Engineering Research and Practice*, pages 601–607. Citeseer, 2004.

[Som04]   Ian Sommerville. Software engineering. international computer science series. *ed: Addison Wesley*, 2004.

[Ste05]   Friedrich Steimann. Domain models are aspect free. In *Model Driven Engineering Languages and Systems*, pages 171–185. Springer, 2005.

[Sul02]   Gregory T Sullivan. Advanced programming language features for executable design patterns" better patterns through reflection. 2002.

[TBG03]   Tom Tourwé, Johan Brichau, and Kris Gybels. On the existence of the aosd-evolution paradox. *SPLAT: Software engineering Properties of Languages for Aspect Technologies*, 2003.

[TOHSJ99]   Peri Tarr, Harold Ossher, William Harrison, and Stanley M Sutton Jr. N degrees of separation: multi-dimensional separation of concerns. In *Proceedings of the 21st international conference on Software engineering*, pages 107–119. ACM, 1999.

[YJ02]   Joseph W Yoder and Ralph Johnson. The adaptive object-model architectural style. In *Software Architecture*, pages 3–27. Springer, 2002.

Appendix A

Schema Loading

A.1 Load method

```java
public static Schema load(
  SchemaFactory factory, Class<?>... schemaKlassesDef) {

  // Filter out primitives by loading them separately
  final List<Class<?>> schemaKlasses = new LinkedList<>();
  for (Class<?> schemaClass : schemaKlassesDef) {
    if (Primitives.class.isAssignableFrom(schemaClass)) {
      primitiveManager.loadPrimitives(schemaClass);
    } else {
      schemaKlasses.add(schemaClass);
    }
  }

  // create an empty schema using the factory, will wire it later
  final Schema schema = factory.Schema();

  // build the types from the schema klasses definition
  final Set<Type> types = buildTypesFromClasses(factory, schema, schemaKlasses);

  // wire the types on schema
  // it is inverse so it will refer to schema.types() directly
  types.forEach(type -> type.schema(schema));

  // get the schema's schemaKlass
  final Klass schemaSchemaKlass = factory.Klass().schemaKlass();

  // wire the schema's schemaKlass
  schema.schemaKlass(schemaSchemaKlass);
  return schema;
}
```

Listing A.36: SchemaLoader load method

A.2 Build Types Method

```java
private static Set<Type> buildTypesFromClasses(
  SchemaFactory factory,
  Schema schema,
  List<Class<?>> schemaKlassesDefinition) {
  Map<Type, TypeWithClass> types = new LinkedHashMap<>();

  // <classNameFieldNameCombo, FieldWithMethod>
  Map<String, FieldWithMethod> allFieldsWithReturnType = new LinkedHashMap<>();

  // Klasses
  for (Class<?> schemaKlassDefinition : schemaKlassesDefinition) {
    String klassName = schemaKlassDefinition.getSimpleName();

    Map<String, Field> fieldsForKlass =
      buildFieldsFromMethods(
        klassName, factory, schemaKlassDefinition, allFieldsWithReturnType);

    // create a new klass
    Klass klass = factory.Klass();
    klass.name(klassName);
    klass.schema(schema);

    // wire the owner klass in fields, it is inverse refering to klass.fields()
    fieldsForKlass.values().forEach(field -> field.owner(klass));

    typesCache.put(klass.name(), klass);

    // add the a new klass
    types.put(klass, new TypeWithClass(klass, schemaKlassDefinition));
  }

  // wiring
  wireFieldTypes(factory, schema, allFieldsWithReturnType);
  wireFieldInverse(allFieldsWithReturnType);
  wireFieldTypeKeys(types);

  wireSchemaKlasses(schema.schemaKlass().schema());

  wireKlassSupers(types, typesCache);
  wireKlassSubs(types, typesCache);
  wireKlassClassOf(types, schemaKlassesDefinition);

  typesCache.clear();

  return types.keySet();
}
```

Listing A.37: SchemaLoader buildTypesFromClasses method

A.3   Build Fields Method

```java
public static Map<String, Field> buildFieldsFromMethods(
  String klassName,
  SchemaFactory factory,
  Class<?> schemaKlassDefinition,
  Map<String, FieldWithMethod> allFieldsWithReturnType) {
  Map<String, Field> fieldsForKlass = new LinkedHashMap<>();

  final Method[] fields = schemaKlassDefinition.getMethods();
  // sort methods
  ...

  for (Method schemaKlassField : fields) {
    final String fieldName = schemaKlassField.getName();
    final Class<?> fieldReturnClass = schemaKlassField.getReturnType();

    // skip default method declarations
    if (schemaKlassField.isDefault()) {
      continue;
    }

    // check for many
    final boolean many = primitiveManager.isMany(fieldReturnClass);

    // check for optional
    final boolean optional = schemaKlassField.isAnnotationPresent(Optional.class);

    // check for key
    final boolean key = schemaKlassField.isAnnotationPresent(Key.class);

    // check for contain
    final boolean contain = schemaKlassField.isAnnotationPresent(Contain.class);

    // add its fields, the owner Klass will be added later
    final Field field = factory.Field();
    field.name(fieldName);
    field.contain(contain);
    field.key(key);
    field.many(many);
    field.optional(optional);

    fieldsForKlass.put(fieldName, field);

    // use klassName and fieldName combo here,
    // because the real hashCode can not be calculated.
    allFieldsWithReturnType.put(
      klassName + fieldName, new FieldWithMethod(field, schemaKlassField));
  }
  return fieldsForKlass;
}
```

Listing A.38: SchemaLoader buildFieldsFromMethods method

A.4 Wire Types Method

```java
private static void wireFieldTypes(
  SchemaFactory factory,
  Schema schema,
  Map<String, FieldWithMethod> allFieldsWithReturnType) {

  for (String klassNameFieldNameCombo : allFieldsWithReturnType.keySet()) {
    Method method = allFieldsWithReturnType.get(klassNameFieldNameCombo).method;
    Field field = allFieldsWithReturnType.get(klassNameFieldNameCombo).field;

    // In case the field is multi value (many), that means that the real type is
    // not given in the method.getReturnType() because this will give Set ot List,
    // BUT the real type is in method.getGenericReturnType().
    Class<?> fieldTypeClass;

    // in case it is multi field, get the return the Generic Return Type
    if (field.many()) {
      // The type in this case will be Set or List,
      // but the Generic Return Type will be the actual type.
      ParameterizedType fieldManyType =
        (ParameterizedType) method.getGenericReturnType();
      fieldTypeClass = (Class<?>) fieldManyType.getActualTypeArguments()[0];
    } else {
      fieldTypeClass = method.getReturnType();
    }

    Type fieldType = getFieldType(fieldTypeClass, schema, factory);
    field.type(fieldType);
  }
}
```

Listing A.39: SchemaLoader wireFieldTypes method

Listing A.36 demonstrates the loading process in managed data. As illustrated we first implement the instances and following that we use setters to wire them up. The reason for this is that not everything exists at the time that it needs to be set. Listing A.37 show how the klass creation is performed during the schema loading. Listing A.38 illustrates how the fields of a klass are created reflection. Listing A.39 presents the way the types are wired during the schema creation.

This implementation shows the usage of Java reflection in our implementation. However, because Java reflection capabilities are limited, this restricted our implementation.

Appendix B

The MObject class

This class is the managed data interpreter; therefore, it first it setups the fields of a managed object based on the schema klass. Next, it handles all the calls to the managed objects' methods and finally is responsible for type checking of the fields' values during initialization or assignment.

The usage of the `schemaKlass` for setting up the fields is shown in Listing B.40.

```java
public MObject(Klass schemaKlass, Object... initializers) {
  this.schemaKlass = schemaKlass;
  this.schemaKlass.fields().forEach(this::safeSetupField);
  this.safeInitializeProps(initializers);
}

protected void setupField(Field field) {
  if (!field.many()) {
    if (field.type().schemaKlass().name().equals("Primitive")) {
      this.props.put(field.name(), new MObjectFieldSinglePrimitive(this, field));
    } else {
      this.props.put(field.name(), new MObjectFieldSingleMObj(this, field));
    }
  } else {
    if (field.type().schemaKlass().name().equals("Primitive")) {
      this.props.put(field.name(), new MObjectFieldManyList(this, field));
    } else {

      Klass klassType = (Klass) field.type();
      if (klassType.key()!= null) {
        this.props.put(field.name(), new MObjectFieldManySet(this, field));
      } else {
        this.props.put(field.name(), new MObjectFieldManyList(this, field));
      }
    }
  }
}
```

Listing B.40: MObject: setup fields

The `schemaKlass` is given to the `MObject` by the `DataManager` that is responsible for creating it. Using this `schemaKlass` the `MObject` setups the `Fields` of the Klass, Line 3. Inside the `setupField` method the interpretation of the schema is performed. In particular, in Line 8 we check if that field is a multi-value field, and if not, we just set it up as a `Primitive` or a `Klass` accordingly. Consider that the `field.type().schemaKlass().name()` is used like a common `instanceof` in Line 9. In case the

field has many values, we first check if it is `Primitive`, since we do not support Set of `Primitive`s. Following that, we check if a `Key` field exists on that field's type and in that case this field is a Set, otherwise it is a List.

The invocation handling process of managed object is showed in Listing B.41.

```java
public Object invoke(
    Object proxy, Method method, Object[] args) throws Throwable {
    final String fieldName = method.getName();

    if (method.isDefault()) { // if the method is default, invoke this one
        return _callDefaultMethod(proxy, method, args);
    }

    // This is a way to execute the "attached" methods of the derived Managed Objects,
    for (Method declaredMethod : this.getClass().getMethods()) {
        if (declaredMethod.getName().equals(fieldName)) {
            return method.invoke(this, args);
        }
    }

    // Managed Object
    MObjectField mObjectField = this.props.get(fieldName);
    boolean isMany = mObjectField.getField().many();

    if (args == null) {
        return _get(fieldName); // return the field's value
    }

    boolean isAssignment = false;
    Object fieldArgs = args[0];

    if (fieldArgs.getClass().isArray() && ((Object [])fieldArgs).length > 0) {
        isAssignment = true;
    }

    if (isAssignment) {
        if (((Object [])fieldArgs).length == 1 && !isMany) {
            _set(fieldName, ((Object [])fieldArgs)[0]);
        } else {
            _set(fieldName, fieldArgs);
        }
        return null;
    }
    return _get(fieldName);
}
```

Listing B.41: MObject: invocation handler

The type checking for each field is performed by the classes `MObjectFieldSinglePrimitive`, `MObjectFieldSingleMObj`, `MObjectFieldManyList` and `MObjectFieldManySet`.

The basic structure of such class is given in Listing B.42.

```
1   // A field of managed data
2   public abstract class MObjectField {
3
4     // the owner of the field as an Managed Object.
5     protected final MObject owner;
6
7     // the Field.
8     protected final Field field;
9
10    // the Inverse of the field.
11    protected final Field inverse;
12
13    public MObjectField(MObject owner, Field field) {
14      this.owner = owner;
15      this.field = field;
16      this.inverse = field.inverse();
17    }
18
19    // Initializes the field with a value
20    public abstract void init(Object value)
21      throws InvalidFieldValueException, NoKeyFieldException;
22
23    // Checks the given value if it is valid
24    protected abstract void check(Object value)
25      throws InvalidFieldValueException;
26
27    // Returns a default value for this kind of field.
28    protected abstract Object defaultValue()
29      throws UnknownTypeException;
30
31    // Sets a value to the field.
32    public abstract void set(Object value)
33      throws InvalidFieldValueException, NoKeyFieldException;
34
35    // Returns the value of the field
36    public abstract Object get();
37
38    // Returns the Field object that is wrapped.
39    public Field getField() {
40      return this.field;
41    }
42  }
```

Listing B.42: MObjectField abstract class

Appendix C

JHotDraw Migration Process

C.1 DrawingView

One of the main components of JHotDraw is the *DrawingView* interface. As Figure C.1 illustrates, the *DrawingView* is responsible for rendering `Drawings` and listening to its changes. Additionally, it is responsible for receiving the user input and delegating it to the current tool.


Figure C.1: DrawingView of JHotDraw

Conclusively, `DrawingView` makes a good candidate for managed data migration. The reason is that the specifications of that class can be implemented in data managers and dynamically added to it.

C.2 Managed Data DrawingView

To support sub-typing on the `DrawingView` interface, we have implemented the `MDDrawingView`, namely Managed Data DrawingView, which replaced the `DrawingView` in JHotDraw. Having this interface for super type, we still needed the actual managed data schemas. As Figure C.1 shows, there are two implementations of the `DrawingView`. In particular, the `StandardDrawingView`, which is the implementation that is used when a new drawing view is created in the application and the `NullDrawingView`, which represents a null drawing view as for the *null-object* pattern.

Following their original design, we have implemented two schemas, one for the `StandardDrawingView` and one for the `NullDrawingView`, namely `MDStandardDrawingView` and `MDNullDrawingView` respectively. The instances of those schemas have been used in the same way their counterparts are used in JHotDraw. A snippet of the `MDStandardDrawingView` is shown in Listing C.43 [1].

---

[1] Most of the implementation has been omitted for brevity.

```
1  public interface MDStandardDrawingView extends M, MDDrawingView {
2    ...
3    // Composition over inheritance, the original inherits the JPanel
4    JPanel panel(JPanel... panel);
5
6    default JPanel getPanel() {
7        return panel();
8    }
9
10   default void setPanel(JPanel _panel) {
11       panel(_panel);
12   }
13   ...
14   Rectangle damage(Rectangle... damage);
15   Drawing drawing(Drawing... drawing);
16   ...
17   default FigureEnumeration selectionZOrdered() {
18     List result = CollectionsFactory.current().createList(selectionCount());
19     FigureEnumeration figures = drawing().figures();
20
21     while (figures.hasNextFigure()) {
22       Figure f= figures.nextFigure();
23       if (isFigureSelected(f)) {
24         result.add(f);
25       }
26     }
27     return new ReverseFigureEnumerator(result);
28   }
29   ...
30   default void repairDamage() {
31     if (getDamage() != null) {
32       panel().repaint(damage().x, damage().y, damage().width, damage().height);
33       setDamage(null);
34     }
35   }
36   ...
37 }
```

Listing C.43: MDStandardDrawingView schema

Listing C.43 shows that the `MDStandardDrawingView` interface extends both `M` interface, defining that this is a schema definition, and `MDDrawingView`, for sub-type support. Additionally, all the functionalities implemented in methods of the original `DrawingView`, in managed data they are implemented in default methods of the schema interface. The fields of a schema can provide those methods with the managed object's current state. As Lines 17 and 30 show, the fields of the schema can be used to query their values inside the default methods. Note that the code in the default methods is identical to the original `DrawingView`. Furthermore, for consistency with the legacy code, we have implemented setters and getters, Lines 10 and 6, for field values accessors. This way we maintained consistency across in accessing values of the managed object.

A notable issue is that the original `StandardDrawingView` extends the `javax.swing.jpanel` class as Figure C.1 shows. However, such a structure is not supported in managed data. Schema definitions can not extend classes. To overcome this issue we defined the `JPanel` as a field in the schema, namely *panel*. To support the `JPanel` as a type of a field though, it is needs ti be defined as managed data. By all means, the same holds for the remaining fields, such as `Rectangle` and `Drawing`.

As explained in Section 4.4.2, our framework provides external primitives definition by inheriting the `Primitives` interface. The JHotDraw primitives definition is shown in Listing C.44.

```
public interface JHotDrawPrimitives extends Primitives {
  javax.swing.JPanel JPanel();

  java.awt.Color Color();
  java.awt.Cursor Cursor();
  java.awt.Point Point();
  java.awt.Dimension Dimension();
  java.awt.Rectangle Rectangle();

  CH.ifa.draw.framework.DrawingEditor DrawingEditor();
  CH.ifa.draw.framework.Drawing Drawing();
  CH.ifa.draw.framework.Painter Painter();
  CH.ifa.draw.framework.PointConstrainer PointConstrainer();


  ...
}
```

Listing C.44: JHotDraw Primitives Definition

This has been proven very helpful since we did not need to re-implement every field as managed data during the refactoring. Especially, classes that are provided by libraries such as `javax.swing` and `java.awt`.

C.3   MDDrawingView Schema Factories

In order to create instances of the defined `MDStandardDrawingView` and `MDNullDrawingView` schemas, we needed their factories. Besides the schema factories, which is as simple as Listing C.45 shows, we still needed a way to give initialization values to the schema instances the same way that the original `StandardDrawingView` does during construction. Additionally, this factory should be used like Java's `new` keyword in the source code. This factory just replicates the original `StandardDrawingView` constructor and is used from the program to create new instances of the schemas. The code of the `MDStandardDrawingView` *instances factory* is illustrated in Listing C.3, in comparison to the original constructor, illustrated in Listing C.3.

```
public interface DrawingViewSchemaFactory extends IFactory {
  MDStandardDrawingView DrawingView();
  MDNullDrawingView NullDrawingView();
}
```

Listing C.45: DrawingView Schema Factory

```
1  public StandardDrawingView(DrawingEditor editor, int width, int height) {
2    setAutoscrolls(true);
3    fEditor = editor;
4    fViewSize = new Dimension(width,height);
5    setSize(width, height);
6    fSelectionListeners = CollectionsFactory.current().createList();
7    addFigureSelectionListener(editor());
8    setLastClick(new Point(0, 0));
9    fConstrainer = null;
10   fSelection = CollectionsFactory.current().createList();
11   setDisplayUpdate(createDisplayUpdate());
12   setBackground(Color.lightGray);
13   addMouseListener(createMouseListener());
14   addMouseMotionListener(createMouseMotionListener());
15   addKeyListener(createKeyListener());
16 }
```

Listing C.46: Original StandardDrawingView Constructor

```
1  public static MDDrawingView newDrawingView(
2    DrawingEditor editor, int width, int height) {
3    final MDStandardDrawingView drawingView = drawingViewSchemaFactory.DrawingView();
4    MyJPanel jPanel = new MyJPanel();
5    jPanel.setAutoscrolls(true);
6    jPanel.setSize(width, height);
7    jPanel.setBackground(Color.lightGray);
8    drawingView.panel(jPanel);
9    jPanel.setDrawingView(drawingView);
10
11   drawingView.editor(editor);
12   drawingView.size(new Dimension(width, height));
13   jPanel.setSize(width, height);
14   drawingView.lastClick(new Point(0, 0));
15   drawingView.constrainer(null);
16   drawingView.setDisplayUpdate(new SimpleUpdateStrategy());
17   drawingView.setBackground(Color.lightGray);
18   drawingView.drawing(new StandardDrawing());
19
20   jPanel.addMouseListener(...);
21   jPanel.addMouseMotionListener(...);
22   jPanel.addKeyListener(...);
23   return drawingView;
24 }
```

Listing C.47: MDStandardDrawingView Instances Factory

C.4 MDDrawingView Integration

Finally, in order to integrate the `MDDrawingView` managed objects in the existing system, first we had to replace every instance of `DrawingView` with `MDDrawingView`, every `StandardDrawingView` with `MDStandardDrawingView` and every `NullDrawingView` with `MDNullDrawingView` accordingly. In addition, everywhere a new instance of these is created, we replaced it with our *instances factory*.

For instance Listings C.1 and C.2 show how the code has been changed in the `DrawApplication` class.

```
1  Dimension d = getDrawingViewSize();
2
3  DrawingView newDrawingView =
4    new StandardDrawingView(this, d.width, d.height);
5
6  newDrawingView.setDrawing(newDrawing);
```

Listing C.1: Original createDrawingView

```
1  Dimension d = getDrawingViewSize();
2
3  MDDrawingView newDrawingView =
4    MDDrawingViewFactory
5      .newSubjectRoleDrawingView(this, d.width, d.height);
6
7  newDrawingView.setDrawing(newDrawing);
```

Listing C.2: ManagedData createDrawingView

The factory's code of the `MDDrawingView` with the observable data manager is shown in Listing C.48.

```
1  public static MDDrawingView newSubjectRoleDrawingView(
2    DrawingEditor editor, int width, int height) {
3
4    Schema drawingViewSchema = SchemaLoader.load(
5        schemaFactory,
6        JHotDrawPrimitives.class, MDStandardDrawingView.class);
7
8    FigureSelectionListenerSubjectRoleDataManager subjectRoleFactory =
9      new FigureSelectionListenerSubjectRoleDataManager();
10
11   DrawingViewSchemaFactory drawingViewSchemaFactory = subjectRoleFactory.factory(
12       DrawingViewSchemaFactory.class, drawingViewSchema);
13
14   MDStandardDrawingView drawingView = drawingViewSchemaFactory.DrawingView();
15
16   MyJPanel jPanel = new MyJPanel();
17   jPanel.setAutoscrolls(true);
18   ....
19   drawingView.editor(editor);
20
21   drawingView.size(new Dimension(width, height));
22   jPanel.setSize(width, height);
23
24   ...
25
26   // Panel events
27   jPanel.addMouseListener(new MouseAdapter() {...});
28   jPanel.addMouseMotionListener(new MouseMotionListener() {...});
29   jPanel.addKeyListener(new DrawingViewKeyListener(drawingView));
30
31   return drawingView;
32 }
```

Listing C.48: ManagedDataJHotDraw: MDDrawingView Factory

Appendix D

Metrics Results

D.1 JHotDraw Results

D.1.1 FigureSelectionListener

	Coupling		Cohesion		Size		Separation of Concerns		
Class	CBC	DIT	LCOO	WOC	LOC	NOA	CDC	CDO	CDLOC
DrawingView (Subject)	-	-	-	-	54	-	1	2	-
StandardDrawingView (concrete Subject)	35	5	17	136	629	21	1	9	20
NullDrawingView (concrete Subject)	18	5	46	54	155	5	1	2	2
FigureSelectionListener (Observer)	-	-	-	-	4	-	1	1	-
DrawingEditor (Interface)	-	-	-	-	14	-	1	1	-
DrawApplication (concrete Observer)	71	6	2	144	726	20	1	1	4
AbstractCommand (concrete Observer)	33	1	6	28	133	4	1	3	8
UndoableCommand (concrete Observer)	13	1	4	21	78	3	1	2	12

Table D.1: JHotDraw FigureSelectionListener Metrics

	Coupling		Cohesion		Size			Separation of Concerns		
	CBC	DIT	LCOO	WOC	LOC	NOA	VS	CDC	CDO	CDLOC
Total	170	18	75	383	1793	53	8	8	21	46
Max	71	6	46	144	726	21	-	-	9	20
Min	13	1	2	21	4	3	-	-	1	2
Average	34	3.6	15	76.6	224.125	10.6	-	-	2.625	9.2

Table D.2: JHotDraw FigureSelectionListener Totals

D.1.2 ChangeAttributeCommand

	Coupling		Cohesion		Size		Separation of Concerns		
Class	CBC	DIT	LCOO	WOC	LOC	NOA	CDC	CDO	CDLOC
Command	-	-	-	-	12	-	1	2	-
AbstractCommand	33	1	6	28	133	4	1	4	12
UndoableCommand	13	1	4	21	78	3	1	2	4
ChangeAttributeCommand	11	2	2	5	103	2	1	3	6
DrawApplication	71	6	2	144	726	20	1	5	10

Table D.3: JHotDraw ChangeAttributeCommand Metrics

	Coupling		Cohesion		Size			Separation of Concerns		
	CBC	DIT	LCOO	WOC	LOC	NOA	VS	CDC	CDO	CDLOC
Total	128	10	14	198	1052	29	5	5	16	32
Max	71	6	6	144	726	20	-	-	5	12
Min	11	1	2	5	12	2	-	-	2	4
Average	32	2.5	3.5	49.5	210.4	7.25	-	-	3.2	8

Table D.4: JHotDraw ChangeAttributeCommand Totals

D.2 ManagedDataJHotDraw Results

D.2.1 FigureSelectionListener

	Coupling		Cohesion	Size			Separation of Concerns		
Class / Data manager	CBC	DIT	LCOO	WOC	LOC	NOA	CDC	CDO	CDLOC
MDDrawingView (Interface)	-	-	-	-	111	-	0	0	0
MDStandardDrawingView (Managed Data)	26	2	16	120	471	18	0	0	0
MDNullDrawingView (Managed Data)	15	2	31	38	120	8	0	0	0
MDDrawingViewFactory (Helper)	17	1	1	8	177	-	1	1	0
DrawingViewSchemaFactory (SchemaFactory)	-	-	-	-	4	0	0	0	0
FigureSelectionListenerMObject (Data Manager)	5	3	4	6	28	0	1	4	0
DrawingEditor (Interface)	-	-	-	-	12	-	0	0	-
DrawApplication (concrete Observer)	72	6	2	145	732	20	1	1	2
AbstractCommand (concrete Observer)	33	1	6	28	133	4	0	0	0
UndoableCommand (concrete Observer)	16	1	4	23	90	3	1	2	4

Table D.5: MDJHotDraw FigureSelectionListener Metrics

	Coupling		Cohesion		Size		Separation of Concerns		
	CBC	DIT	LCOO	WOC	LOC	NOA	CDC	CDO	CDLOC
Total	184	16	64	368	1878	53	4	8	6
Max	72	6	31	145	732	20	-	4	4
Min	5	1	1	6	4	0	-	0	0
Average	26.285	2.285	9.142	52.571	187.8	7.571	-	0.8	0.66

Table D.6: MDJHotDraw FigureSelectionListener Totals

D.2.2 ChangeAttributeCommand

	Coupling		Cohesion		Size			Separation of Concerns		
Class / Data manager	CBC	DIT	LCOO	WOC	LOC	NOA	VS	CDC	CDO	CDLOC
Command	-	-	-	-	10	-	1	0	-	0
AbstractCommand	33	1	6	28	133	4	1	0	0	0
UndoableCommand	16	1	3	23	90	3	1	1	1	2
UndoableChangeAttrCmdMObject	7	3	1	3	98	0	1	1	2	0
DrawApplication	72	6	2	145	732	20	1	1	5	10

Table D.7: MDJHotDraw ChangeAttributeCommand Metrics

	Coupling		Cohesion		Size			Separation of Concerns		
	CBC	**DIT**	**LCOO**	**WOC**	**LOC**	**NOA**	**VS**	**CDC**	**CDO**	**CDLOC**
Total	128	11	12	199	1063	27	5	3	8	12
Max	72	6	6	145	732	20	-	-	5	10
Min	7	1	1	3	10	0	-	-	0	0
Average	32	2.75	3	49.75	212.6	6.75	-	-	2	2.4

Table D.8: MDJHotDraw ChangeAttributeCommand Totals

D.3 Metrics Comparison Graphs

D.3.1 FigureSelectionListener

Figure D.1: FigureSelectionListener Coupling

Figure D.2: FigureSelectionListener Cohesion

Figure D.3: FigureSelectionListener Size

Figure D.4: FigureSelectionListener Separation of Concerns

### D.3.2 ChangeAttributeCommand

Figure D.5: ChangeAttributeCommand Coupling

Figure D.6: ChangeAttributeCommand Cohesion

Figure D.7: ChangeAttributeCommand Size

Figure D.8: ChangeAttributeCommand Separation of Concerns